

Embedded Real-Time Systems (TI-IRTS)

Project Report for Sapien 190 (PSIMU)
Spring 2010



Peter Høgh Mikkelsen (20087291)

Anders Block Arnfast (20085515)

Kim Bjerger (20097553)

PAsien

Preface

The prerequisite for writing this project report and its accompanying "Architecture Description" document are to pass the graduate course "Embedded Real-Time Systems" (TIIRTS) at the Engineering College of Aarhus. The bases for the project are the exercises described in the handed-in journal: "Exercises 1-5".

The subjects involved: Object Oriented Analysis & Design and patterns for embedded real-time systems are key areas at the Distributed Real-Time Systems specialization at the Engineering College of Aarhus.

The exercise is about designing a product, but the scope of the course is software architecture and design, so this will be the focus of this report, rather than the actual product.

Table of contents

1. Introduction	1
1.1. Teaching OOA/D at IHA	1
1.2. Applying Patterns in Real-Time Embedded Systems	2
1.3. Problem Statement	2
1.4. Presentation of the Patient Simulator System	2
1.5. Report structure	3
2. Project Description (Kim).....	4
2.1. Project execution (Kim).....	4
2.2. Analysis process and methods	10
2.3. Design process and methods.....	15
2.3.1. Deployment view	22
2.3.2. Logical view - discrete package.....	23
2.3.3. Logical view - continuous package.....	27
2.3.4. Logical view - communication package.....	31
2.3.5. Process view	32
2.3.6. Rate Monotonic Analysis	38
2.4. Translation and testing	39
2.5. Development tools.....	42
2.5.1. QT Creator	42
2.5.2. Rhapsody	42
2.5.3. Eclipse.....	42
2.5.4. Subversion and Google Code.....	42
2.6. Results.....	42
2.7. Discussion on achieved results	46
2.8. Experience obtained	48
2.9. Excellence of the project	49
2.10. Suggestion to improvements	50
3. Conclusion	51
4. References.....	53

1. Introduction

What is a Real-Time Embedded System, what characterizes it and how do we create a robust, maintainable- and efficient software design for it?

An embedded system is traditionally characterized as a custom build system made for a special purpose, with limited power, memory and processing resources. A Real-Time System is characterized by interaction with the outside environment and having performance deadlines to meet. An example of a real-time embedded system is a potato sorter: Designed specifically for the food processing industry, it weighs the potatoes while in motion on the wavier and sorts them at the drop-out. This requires special hardware as it must fit in a watertight enclosure where a cooling fan is banned. It also requires real-time processing of weight input to control the sorter mechanism, thus it is an embedded real-time system.

Traditionally these kinds of systems have been built on an 8-bit microcontroller platform with code written in assembler or C. As the system matures, the code base evolves. The code becomes hard to maintain and to extend and competition pressures the development time; the software calls for a change in paradigm.

Several companies¹ have moved their real-time embedded software development to use design patterns and the methods from Object Oriented Analysis and Design. They have hereby gained better products that are easier to maintain and extend. This is the motivation for this course as it is explained next.

1.1. Teaching OOA/D at IHA

OOA/D is used extensively throughout the under-graduate studies of the Information Technology Engineering department at IHA. It is however primarily taught in a non-time critical Windows PC context, rather than in an embedded real time context, which is the scope of this graduate course.

The course is designed to alternate between lecture and practical exercises. The lectures have been split into two sections: "Real-Time Design Patterns" primarily referring to the book by Bruce Powel Douglas (Douglas, 2002) and "Design Patterns – GoF" referring to the book by Gamma et al (Gamma, 1994).

The learning objectives for the course are to analyze and describe requirements for an embedded real-time systems and select and use design patterns for the implementation. An additional objective is to develop product documentation using UML notation.

¹ <http://staff.iha.dk/foh/Foredrag/TwoPartArchitecturalModel-Paper.pdf>

The skills are taught in the lectures, and their application learned through the exercises, in particular in this final hand-in.

1.2. Applying Patterns in Real-Time Embedded Systems

Design patterns originate back to the Small-Talk days of Xerox Parc, where the MVC “pattern” was first invented by Trygve Reenskaug in 1979². It was not until the release of the famous GOF book, that the usage of design patterns became widely-spread.

A design pattern is a named software design that describes the solution to a problem in such a way that you can use it over and over again (Douglas, 2002 p. xiv). Bruce Powel Douglas was with his book in 2002 the first to take design patterns into the real-time embedded domain. Design patterns support the software development at both architectural and mechanistic levels. The real-time patterns are specialized towards resource specific issues, such as memory handling and process scheduling. There is a plethora of design patterns and a range of them will be investigated in this report.

1.3. Problem Statement

The basis for this report is the implementation of a Patient Simulator System. The product implementation is used for application and reflection of design choices. In line with the problems introduced earlier, we state the problem as:

“How do we design an embedded real-time system that fulfils the requirements of the Project Description³ and incorporates design patterns and OOA/D methodology?

What architecture- and design patterns should be incorporated to create a design that is easy to extend and maintain and yet has acceptable real-time performance on the embedded platform?”

The following section will give a short introduction to the Patient Simulator System and set the scope for the implementation according to the problem statement.

1.4. Presentation of the Patient Simulator System

For this exercise we could choose between two products to implement, a Patient Simulator System (PSIMU) or a Local Monitor System (LMON).

The Patient Simulator System project was chosen because it raised a lot of interesting real-time design questions: How do we transform discrete measurements into a real-time output that also acts on

² <http://heim.ifi.uio.no/~trygver/themes/mvc/mvc-index.html>

³ <http://kurser.iha.dk/eit/tiirts/Projekter/PSIMU-project.doc>

external real-time inputs? How can we design a system that is very flexible when it comes to patient scenarios?

The purpose of the PSIMU is to simulate human physiological signals that can be measured by a patient monitor system. The simulated signals are based on real recordings taken from the PhysioBank database. The database signals are post-processed by mathematical models that take inputs from example an injection pump.

Based on the requirement specification and the problem statement, we have chosen to limit the project to focus on the implementation of use case #1. In doing so, we can keep focus on software design, rather than product finalization. Another important restriction in our implementation is that our system only works in soft-real-time. The Patient Simulator is not life support equipment; its usage is only for educational purposes, so missing a deadline will not be catastrophic.

More information on the product documentation can be found in the requirement specification reference [5] and architecture document reference [4].

1.5. Report structure

This report is split into three major sections: Introduction, Project Description and Conclusion. The Project description is the theoretical section where theory is applied to the project. It presents the development model that we have used in our work, the architectural- and design related considerations done during the development process and the results of our analysis and design. The theory is based on the books by Douglas (Douglas, 2002) and Gamma et al (Gamma, 1994) and relevant articles. A complete list can be found in the Reference section and the end of the report.

prototypes with different purpose to investigate the platform, technologies and finally the implementation of the functionality for the product.

In the analysis phase we have started with the requirement analysis by delivering a use case specification⁵ for the Sapien 190 product to our "Customer = Teacher". Here we have identified a number of actors and use cases as shown in Figure 2.

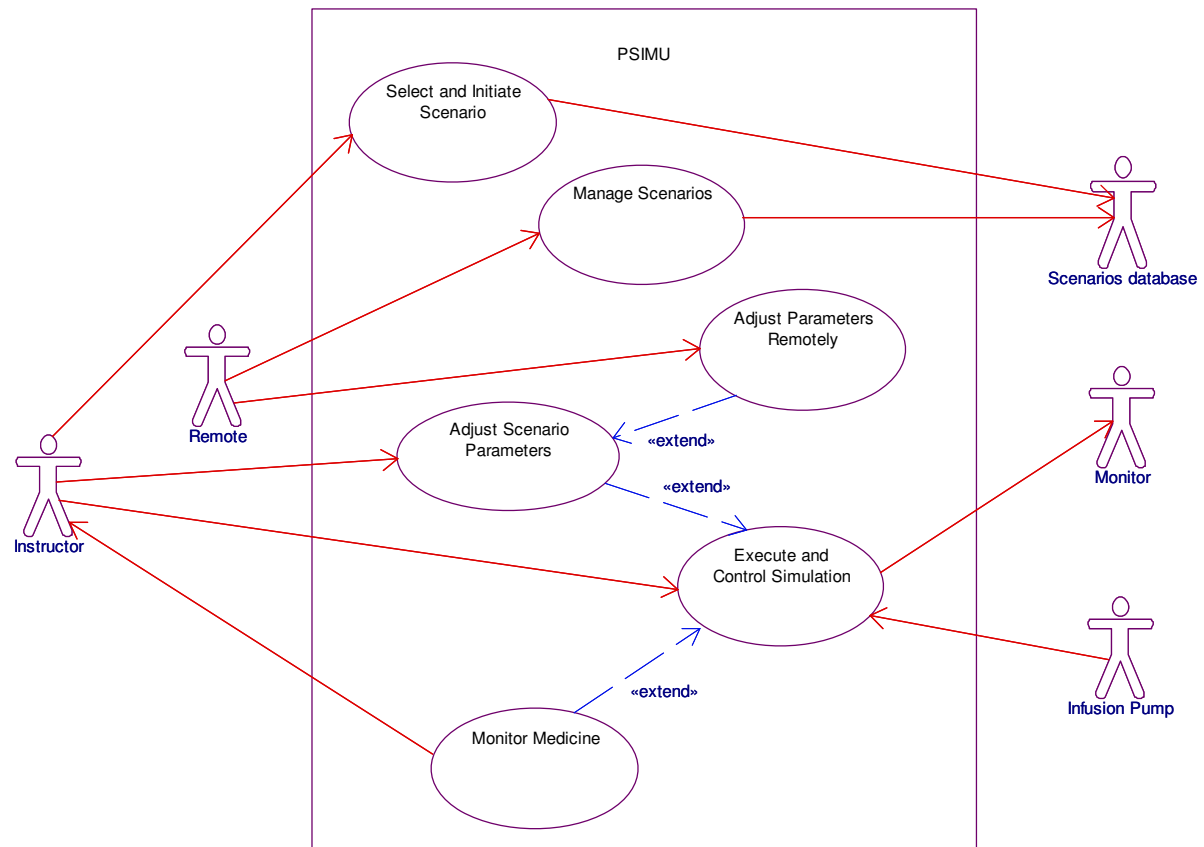


Figure 2 Use case specification and analysis

The next step was to make a domain model for the use case "Execute and Control Simulation" and complete the first iteration in the ROPES spiral micro-cycle. The use case "Execute and Control simulation" is the most complex and significant for the architecture of the real-time part of the patient simulator. This first official prototype we have delivered to our "Customer" together with a first version of the architectural documentation and a demonstration of the Sapien 190 patient simulator. In the W-Model⁶ by Alistair Cockburn he defines external visible deliveries for the external stakeholders of a project being able to follow the progress and give feedback to the project.

⁵ See use case specification for Sapien 190 in references [5]

⁶ Slide 43 from Lesson 8 – L3_ROPES_process

First Delivery 11. May 2010

- Updated Requirement Specification
- Draft Product Architecture Document
- Status Report
- Simulator prototype first version (Part of Use Case #1)

Final Delivery 4. June 2010

- Requirement Specification
- Product Architecture Document
- Project Report
- Simulator prototype (Use Cases parts of #1, #2 and #3)

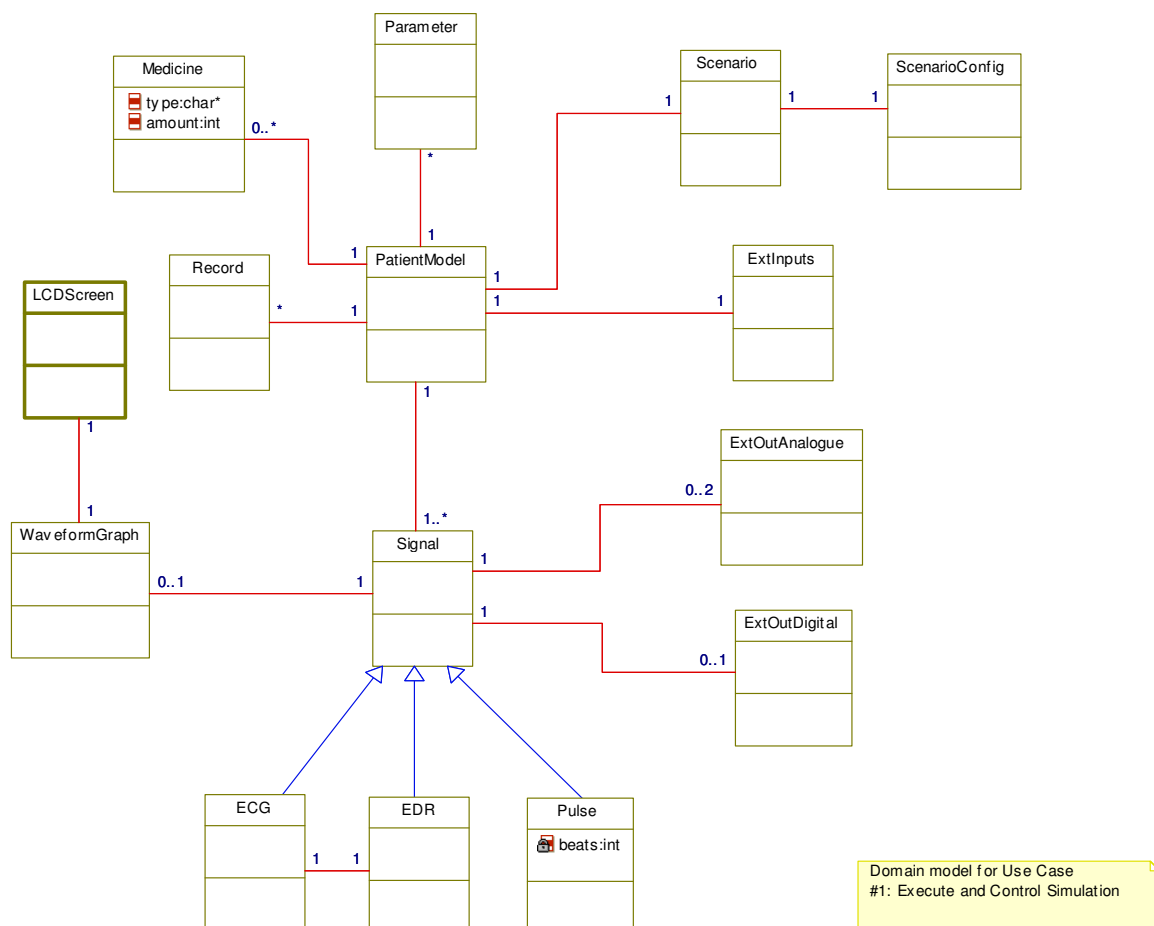


Figure 3 The first domain model for UC#1 Control and Execute Simulation

At the same time using the ROPES micro-cycle spiral to develop the actual product we did make a number of experimental prototypes to reduce the number of unknowns and risk in the project. Some of the questions we had in the beginning were:

- Would it be possible to generate the analogue ECG signal from user space in Linux on the target platform with the sampling rate of 250Hz?
- What would the CPU load be?
- How to plot the waveform ECG signals on the LCD screen using Qt?
- How to cross-compile the patient record WFDB library to the Linux target platform?

The first prototype was to reduce risk by investigate how to use the embedded Linux platform in reading patient record on the target and output the analogue ECG signal. The second prototype was to investigate how to develop with Qt and creating a waveform ECG graph based on the patient record readings.

The following prototypes has been about producing functionality according to the use case specification, where for every iteration parts of a use case or more use cases has been included. These prototypes has been developed by a combination of automatic code-generation from a Rhapsody UML model and source code manual written for the HW and OS abstraction layers. The GUI has been written using Qt from Nokia and integrated with the manual code and code generated from Rhapsody.

ROPES describe the key enabling technologies like visual UML modelling, model execution and mode-code associativity. These key technologies has been possible to realize in this project by use of the IBM Rapsody UML design tool not only for drawing UML diagrams, but also for high level modelling and automatic code generation. This approach has enabled the development process to focus on design instead of implementation. The ROPES key technologies has been a help to make fast iterations between testing new design ideas, by animation of state chats, setting breakpoints in Rhapsody and inspecting variables and states in the model. This higher level modelling approach compared to normally coding, debugging and testing has turned the development process to be more design focused than traditional development. We have used the animated sequence diagrams for test documentation and generating the code directly to Linux and the target platform, which has saved us for time fixing typing errors and trivial debugging and test.

The Rational Unified Process (RUP) has a process workflow that specifies a number of phases: Inception, Elaboration, Construction and finally Transition. In each phase a number of iterations are

perform like in ROPES. In the beginning focus is on the process workflow requirements and business modelling. In this project we have primary been working in the elaboration phase working with 2 major iterations of product release. Since focus for this course in some extend has been design patterns, we have use most the time in the elaboration phase with focus on the architectural and mechanistic design of the project.

Process Workflows

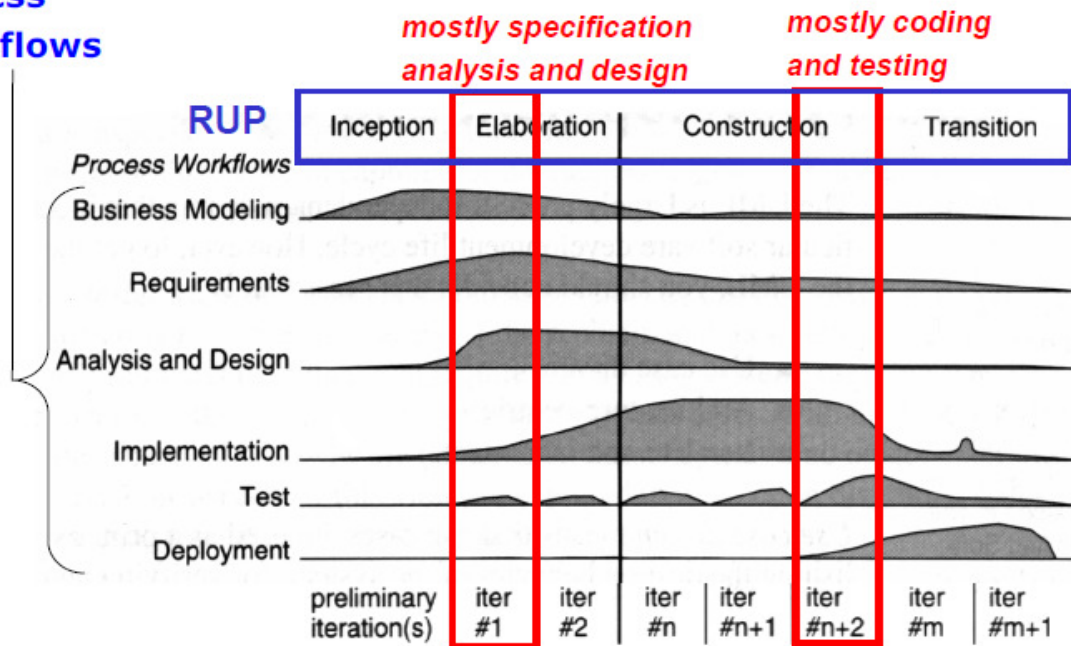


Figure 4 Rational Unified Process workflows

In controlling the progress of the project we have had regularly meetings in which we have been inspired from the way Scrum meetings are organized. In every meeting we have used time on each project member giving a short status of what has been done, issues and problems discovered since last meeting based. The status is based on the planned activities from last meeting. An updated list with notes of meeting status has been updated for every project meetings see appendix A. On every meeting the backlog in Scrum terms has been updated according to our eventually changes priority. We did not define a scrum master or product owner, but on every meeting one of the project members did take the lead on updating the status list and together we did a prioritizing of the backlog (Pending activities). Below see example for the contents of meeting status.

Scrum Status 25. May:

<Name>:

Done:

- Meeting minutes
- Added text to Chapter 5

Problems:

- What exactly to put into ch 11+12

Next:

- Architecture document chapter 11. and 12.

Next scrum meeting Tuesday at 9:00

- Scrum status
- Status on writing Architecture document
- Status on report writing and assignment of tasks
- Continue to 16:00

Action list (Backlog) to final delivery 04. June:

- Finalize Product Architecture Document
- Finalize Project Report
- Implement ECG to Pulse filter

.....

2.2. Analysis process and methods

In the requirement analysis phase with reference to the ROPES spiral model we have created a number of use cases to define the required functionality of the product. A detailed specification of this work can be found in "Requirement Specification for Sapien 190" 5. The first step has been to define the actor-context diagram to identify the actors of the patient simulator. Here we have used the specifications of the PSIMU, LMON, IPUMP and interface specification referenced in [5] chapter 1.2

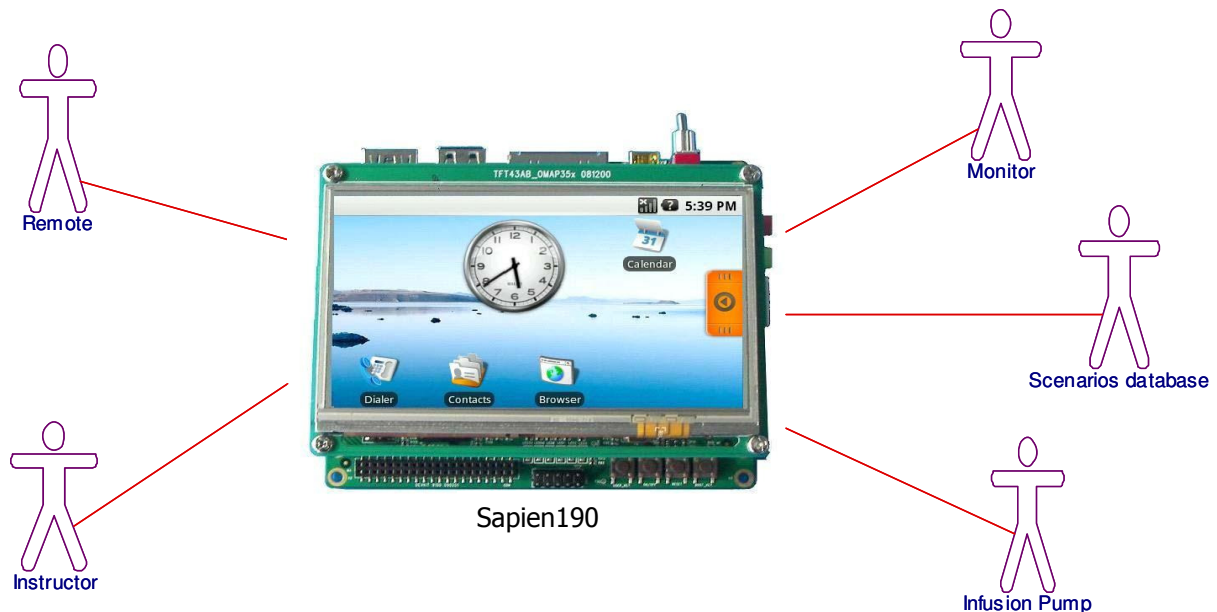


Figure 5 Actor-Context Diagram

The second step in the use case analysis was to identify a number of use cases for each actor as illustrated in use case diagram Figure 2. These use cases are described in details in the identified to be:

1. Execute and Control Simulation
2. Select and Initiate Scenario
3. Adjust Scenario Parameters
4. Monitor Medicine
5. Manage Scenarios

For each use case specification a main scenario and extension is described like on page 9 in the "Requirement Specification for Sapien 190" 5. This document contains the textual description of the system context, interface to external actors, non functional requirements like performance, prioritized product qualities and design constraints. Finally the document contains a description of the original planned deliveries of use cases to be contained in the two first deliveries to our "Customer". We have prioritized maintainability, correctness and usability as the most important quality factors for the

patient simulator. The simulator application is expected to have a longer life time than the hardware and must be able to be maintained for many years in the future. Therefore it is important for the product to focus on these factors in creating the architecture and design.

The domain analysis is based on the use case requirement specification where we did start with the use case "Execute and Control Simulation". This approach is defined in the Unified Process⁷ (UP). A domain model is created to identify the conceptual domain classes and relations between them to give us a better and deeper understanding of the actual problem. We have chosen the "Execute and Control Simulation" use case to be the first in creating this domain model since it is the most complicated and contains the essential functionality for the patient simulator. The UML domain class diagram is illustrated in Figure 3. To this first domain model we added boundary and one control classes. The purpose of the boundary classes is to separate the model from the external actors and the control class is used to encapsulates the control of the scenario described in the UC#1.

⁷ Craig Larman, Applying UML and Patterns, Third Edition chapter 9 Domain Models

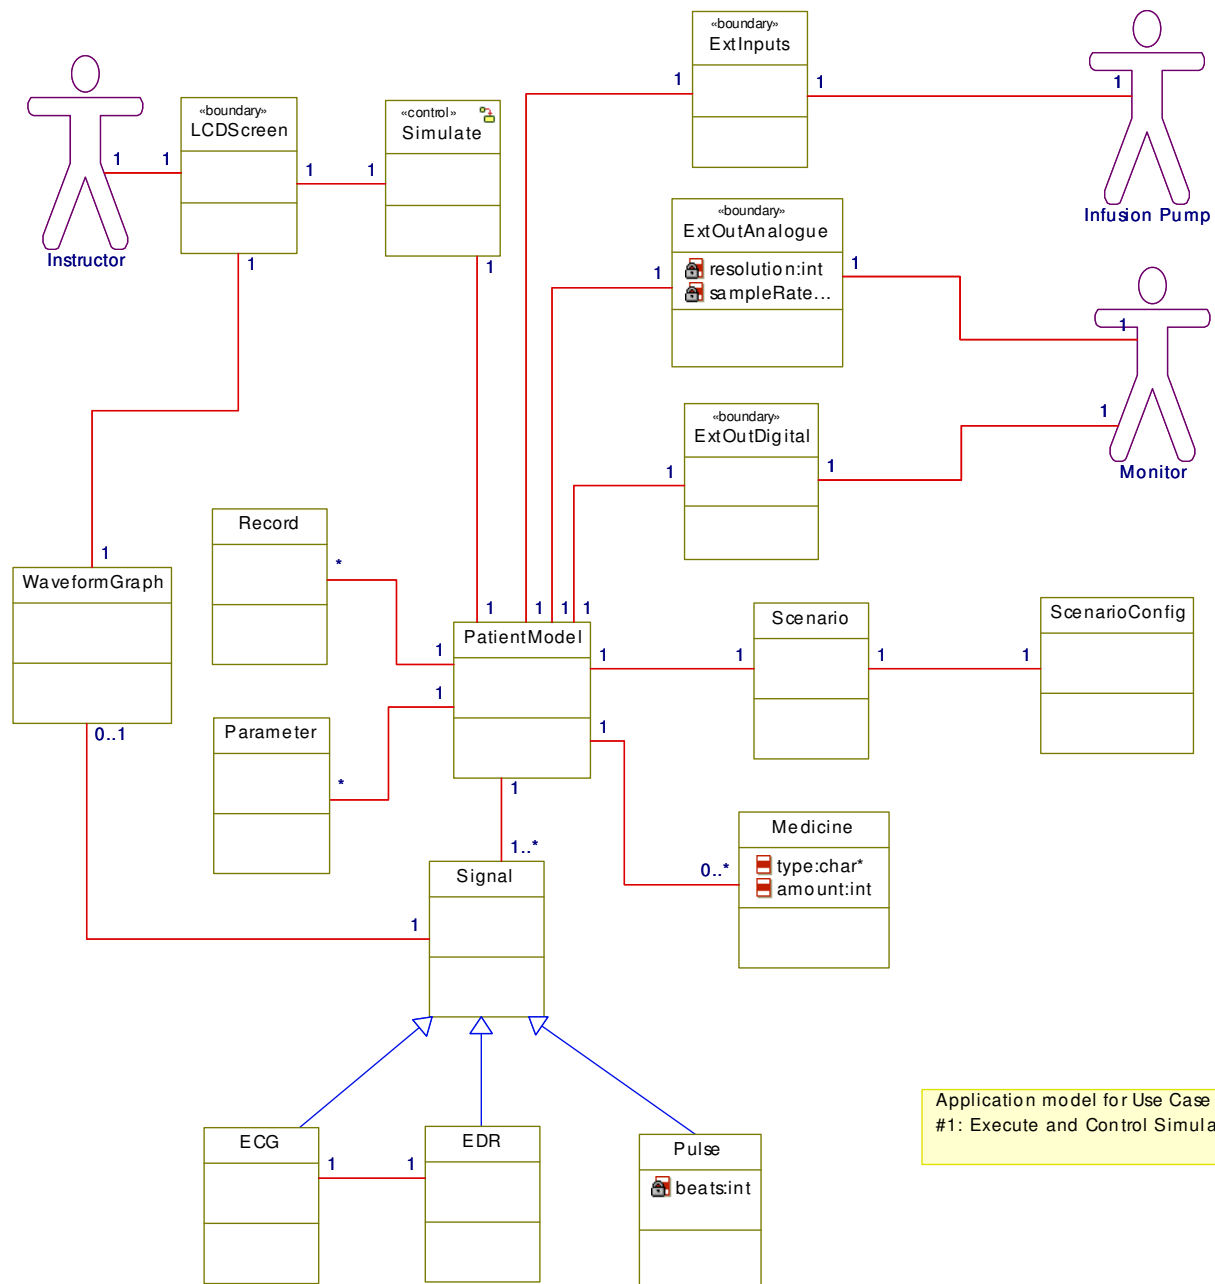


Figure 6 Domain model with boundary and control classes

In the domain model for this first iteration we have identified all the important classes and the relations between them. The domain model is saved in a separate Rhapsody project since the domain model in the following design iterations is changed a number of times. The domain model illustrates how the patient model is the central class for coordinating the patient simulation. When the patient simulation is running it will perform reading of the digitized physiologic signals from a record and send these "real time" values as analogue signals to local connected bedside monitoring equipments. Up to 2 analogue channels with different signals is possible to be simulated simultaneously. The simulated

signals can be ECG or EDR. The pulse will be signalled to the bedside monitoring equipment as a digital signal. The basic idea is that the instructor later should be able to select between different scenarios that describes the configuration of a simulation which include a certain patient model (Normal or with Infusion Pump), patient records from the PhysioBank⁸ database and parameter settings for the simulation like gain and rate for replaying the patient record. More details can be found in the requirement specification 5. In the UML state diagram below we have added a simple functionality for the instructor to start, stop, pause and resume the patient simulation with a fixed scenario configured for the final prototype. That means UC#1, UC#2 and UC#3 is implemented without being able to select a scenario in UC#2.

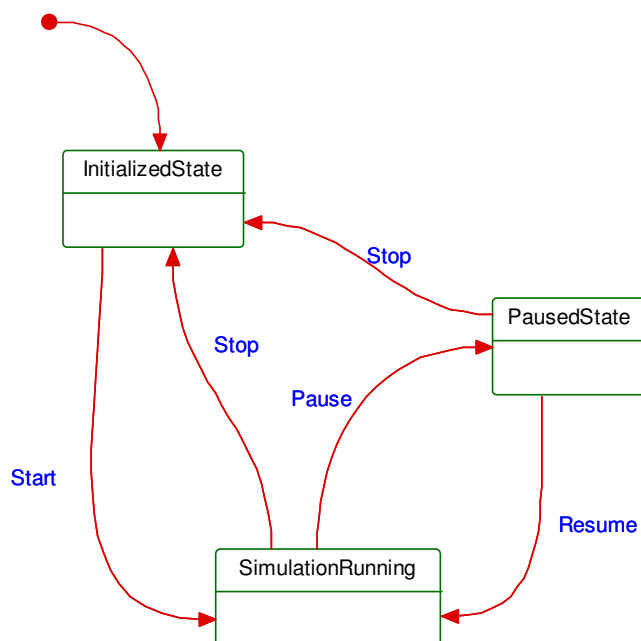


Figure 7 State chart for domain control class simulate

Part of the analysis phase we have identified a number of external and internal generated events that plays an important role in constraining and defining the system behaviour. The events listed below are crucial for analysis of the schedulability and deadlines later in the design process (RMA analysis). For this analysis we have defined number of parameter that can be adjusted to find the optimal timing and scheduling being able to generate the ECG, EDR and the pulse signal. At the same time we should be able to processing information from the IPUMP and updating the waveform signal graph on the LCD screen. A frame buffer has been defined for the number of samples to be collected before updating the LCD screen.

⁸ <http://www.physionet.org/physiobank/physiobank-intro.shtml>

Internal and external event list

#	Event Id	System response	Arrival Pattern	Event Source	Response time
1	Sample	Calculate and generate EDR and ECG signals	Frequency of 250 (Fs) max 400	Internal timer	Less than sample period
2	Pulse	Calculate pulse every 200 (Np) samples	Fs/Np	Internal timer	Less than sample period
3	PDU	Updates medicine information	Every second (Fi)	IPUMP	Less than 1/2 second
4	FrameBuffer	Updates signal graph on LCD display	Every 50 (Nf) samples (1/8 of LCD display in pixels)	Internal timer	Less than period updating framebuffer

Parameters identified to be used for RMA analysis in design phase:

Time units		Default value	Units
Fs	Sample frequency	250	Hz
Np	Num samples for pulse calculation	200	Number
Fi	PDU frequency	1	Hz
Nf	Frame buffer size	50	Number

2.3. Design process and methods

We have used the 4+1 view designed by Phillipe Kruchten⁹ as the process driver for the design work in the project. The Use Case view is the driver for focus on what to design in the other 4 views as show in the figure below. We have started with UC#1 defining a number of scenarios that is used as the foundation for the architectural, mechanistic and detailed design that is described in details in the "System/product architecture document for Sapien 190" reference [4].

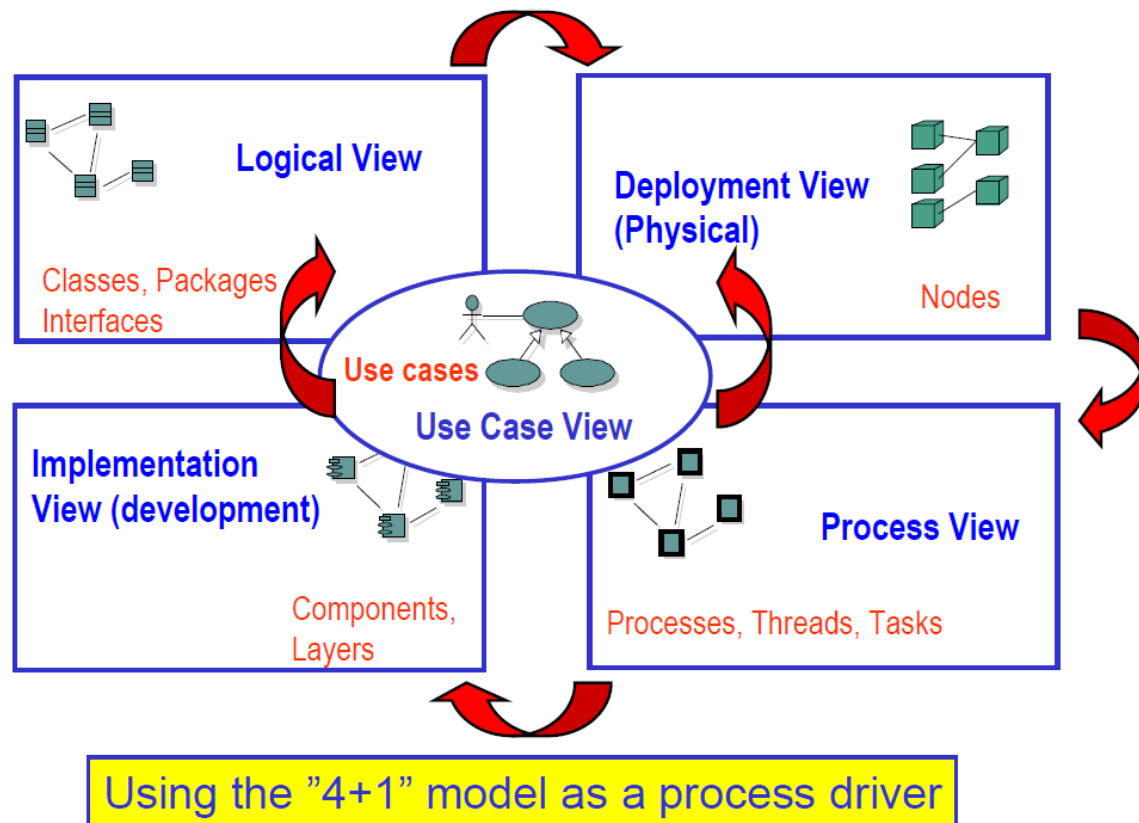


Figure 8¹⁰ "4+1" view model

The UC #1 has been selected for the first iteration of the ROPES spiral microcycle in making the architectural, mechanistic and detailed design. This use case is significant and provides the central functionality of the patient Simulator and contains the real-time constraints. This use case provides the basis functionality that allows the monitor to be connected being able to display the ECG and EDR signals. It also reduces the risk for developing the patient monitor since it covers all the unknown technologies of the product like:

⁹ Kruchten, Phillipe, Architectural Blueprints – The "4+1" View Model of Software Architecture

¹⁰ Slide 4 from Lesson 8, DevelopmentOfRealtimeSystems

- Reading the patient record files on the target (Linux, WFDB and target)
- Generating the analogue output signals (Writing to drivers)
- Display of signal waveform using Qt on target (Working with Qt on target)

The deployment, logical, process and implementation views are designed according to the steps described in development of real-time embedded systems¹¹:

Step 1: Sketch an Actor-Context Diagram

Step 2: Find and document use case #1, #2 and #3 (Chapter 4. and 5.3 Ref. 4)

Step 3: Select a suitable HW architecture (Chapter 7. Ref. 4)

Step 4: Develop a logical model (Chapter 5. Ref. 4)

Step 5: Select a concurrency model (Chapter 6. Ref. 4)

Step 6: Design, Implement, Test and Measure (Chapter 6.5 Rate Monotonic Analysis Ref. 4)

The architectural design is based on the five-layer architecture pattern described in chapter 4.2 reference [2] illustrated in the package diagram shown below.

¹¹ Slides 1 – 17 Lesson 8, DevelopmentOfRealTimeSystem

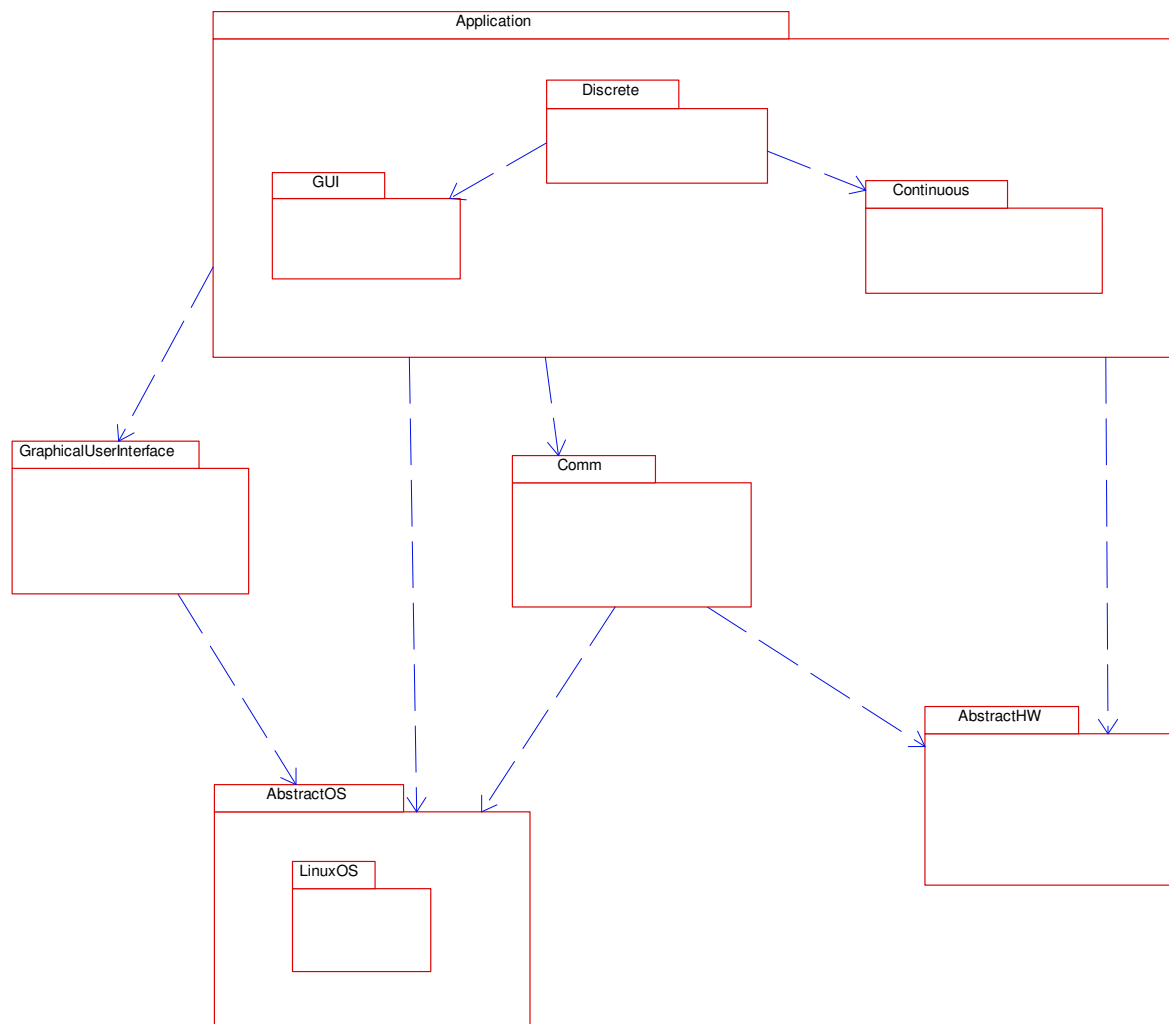


Figure 9 Five layered architecture for logical view

Each abstraction layer is a logical layer representing a well defined domain. Dividing the system into several layers ensures high cohesion for each domain and low coupling between the domains. This simplifies the process of modifying our design or extending it. The application package uses the two-part¹² architectural model for the discrete and continuous parts of the patient Simulator. These packages contain also the most complex part of our design.

The diagram below shows the essential classes defining the architecture for use case #1. One of the important design considerations is to ensure that the dependency between the layers only is in one direction only. Design patterns have been used between the discrete and continuous package (Observer) and between the continuous and communication packages (Mediator) to ensure this one way dependency. Details about use of design patterns will be described in the following sections.

¹² Hans Peter Jepsen, Finn Overgaard Hansen, Designing Event-Controlled Continuous Processing Systems

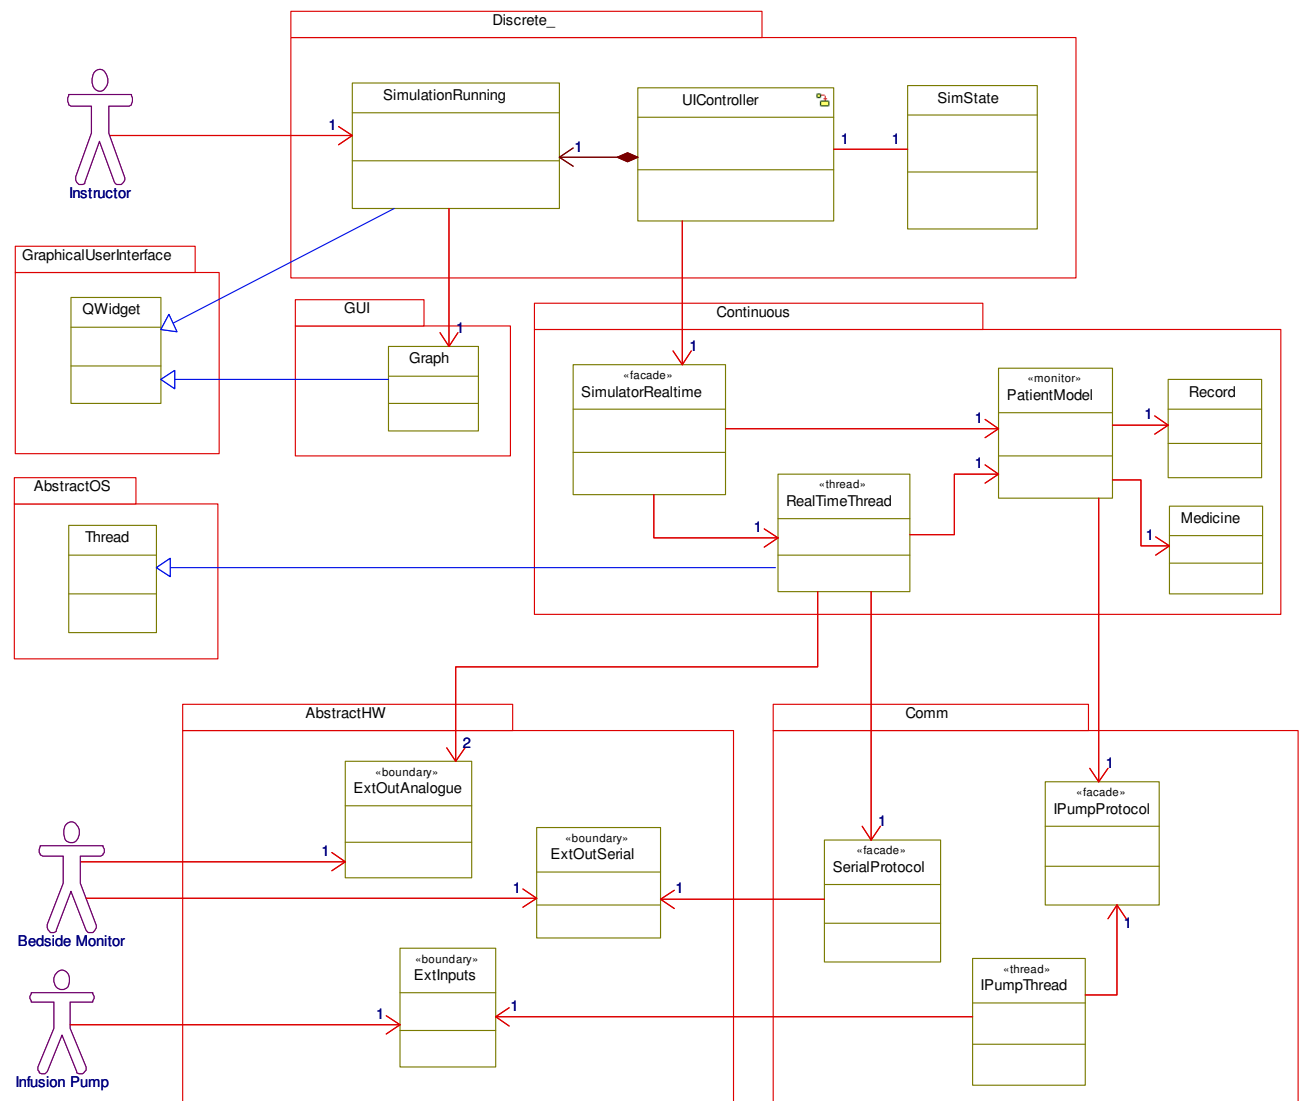


Figure 10 Essential classes in architecture for use case #1 execute and control simulation

After having defined the overall architecture of the design we have focused on the mechanistic design of the continuous and discrete package by refining the domain model in where we have introduced a number of creational, structural and behavioural design patterns from the GoF book (Gamma, 1994).

The following list briefly summarizes the patterns that have been used in the design:

Behavioural Patterns (Gamma, 1994)

- **Observer** (Updating waveform graph) – To notify the UI with a new frame buffer of samples that is updating the waveform graph. By use of this design pattern it would be easy to extend the design with new observers like creation of a FFT view of the ECG signal. This solution separates the continuous package being direct depended on the UI controller in the discrete package.
- **Command** (Setting parameters) – Encapsulates setting parameters in the continuous package by sending a request as an object, thereby letting the UI controller being able to send different type of parameters. It hides the way the parameter is updated in classes that belongs to the continuous package.
- **State** (UI controller in discrete package) – The UI controller depends on its state. The state pattern has been used to handle state changes in combination with the command pattern. Currently this choice is perhaps a bit overkill compared to the simple state machine described in Figure 7, but it will be easy to extend when adding selection of scenarios and new functionality in the future.
- **Command** (In combination with state) – The command pattern has been used in combination with the state pattern to encapsulate generating events to the state machine in the discrete package. This design makes it easy to add new events when the state machine is extended in future software releases.
- **Mediator** (Updating medicine information from IPUMP) – Promotes a loose coupling between the communication package and the continuous package. It provides a way to extend the product by adding different types of devices that can provide inputs to the patient model and thereby manipulating the patient simulation.
- **Strategy** (PhysioModel) – It defines the family of algorithms to compute the patient simulation. Different types of algorithms can easy be added in combination with the filter and pipes pattern. Currently we have designed the NormalModel and InfusionPumpModel more details to be found in chapter 2.3.3.
- **Iterator** (Reading records) – To provide a way to access samples in the records without exposing the underlying representation. The same iterator is used for reading patient records and generating signals for testing.

Structural Patterns (Gamma, 1994)

- **Proxy** – (Reading records) - To provide a place holder for reading records. The proxy pattern has been selected for future extension in where the proxy pattern could be used in combination with the broker pattern as described in chapter 8 reference [2]. This approach would allow us to access records directly from the PhysioBank database on the internet reference [3].

- **Façade** - (Class in continuous package SimulatorRealtime) – To provide a simple unified interface for the subsystems of the real-time part of the patient simulator. It makes it easier for the UI controller to operate on the complex structure of classes in the continuous package.

Creational Patterns (Gamma, 1994)

- **Factory method** (Creating records and filters) – The factory method pattern has been used to make it easy to create new patient record objects. Methods in the façade of the SimulatorRealtime class are provided to generate new objects based on the record file name.
- **Singleton** (DAC + Command State pattern) – Ensure that the states only has one instance and is only created ones. This approach saves a lot of new and deletes operation every time a state is changed in the UI controller. The DAC also uses a modified singleton ensuring we only have one instance per DAC channel. The flyweight¹³ pattern could as an alternative be used to ensure only one object per channel. Would be a better approach moving to a platform with many DAC channels.

Concurrency Patterns (Douglas, 2002)

- **Message Queuing Pattern** (A mailbox of frame buffers) – Used to transfer frame buffers as asynchronous communication messages between the real-time thread and the distributor thread. The frame buffers are used to update the waveform graph by using the observer pattern.
- **Monitor** (Classes PatientModel and FrameBufferPool) – Used to protect shared information between the different threads in the system. The classes PatientModel and FrameBufferPool are designed as monitors since different threads should be synchronized in invocation of methods operating on the same data. This solution separates the communication package being directly dependent on the continuous package.

Memory Patterns (Douglas, 2002)

- **Pool Allocation Pattern** (Allocation of frame buffers) – Creates a pool of frame buffers that is used by the real-time and distributor threads. Frame buffers are created on start-up and available on request by the real-time thread. This approach saves time performing new and delete operation and it ensures a more predictable real-time thread.

Other patterns

- **Filters and pipes** (Filter and calculation of samples) – The filters and pipes pattern is used to compute samples for the different signals (ECG, EDR and Pulse) based on record readings. The structure of the filters and pipes are setup by the strategy (NormalModel and InfusionPumpModel). Different filters are designed for generating the EDR signals, adjusting the

¹³ The Flyweight pattern is a structural pattern in GoF. The pattern is used to share a large number of fine-grained objects efficiently.

gain of the input signal and generating the pulse based on the ECG input signal. The filters and pipes design pattern makes it easy to add more filters or setup of new strategies for computing signals.

In the following chapters we have chosen to give a summary based on the architecture document of some of the essential designs we have implemented. We will only give a summary of the deployment, logical and process view since these views are the most significant for the scope of this project. A fully detailed version of all views is to be found in reference [4]. There will be a discussion on the design considerations and choices we have made. We have also given a more detailed discussion on some of the design patterns described above.

2.3.1. Deployment view

Our target platform is the DevKit8000 that is running Linux. The Sapien 190 application is dependent on Qt libraries and library for the touch screen. The essential hardware components for generating the ECG and EDR signals is the digital to analogue converters (DAC) on the add-on board. Linux drivers were already implemented being able to send one sample value at every write to the driver. A better solution would be to write a more complicated DAC driver that was able to send out a buffer of samples based on a sample rate set from the application. This driver could be implemented using the kernel timer API or creating a process in kernel space. This solution would be more accurate than to use a thread running in user space as we have done in this project. Since Linux driver development is not part of this course we have decided not to choose this path in our development strategy. Actual the solution we have made seems to be good enough as long the target platform is not running other heavy CPU consuming applications like Ethernet network traffic.

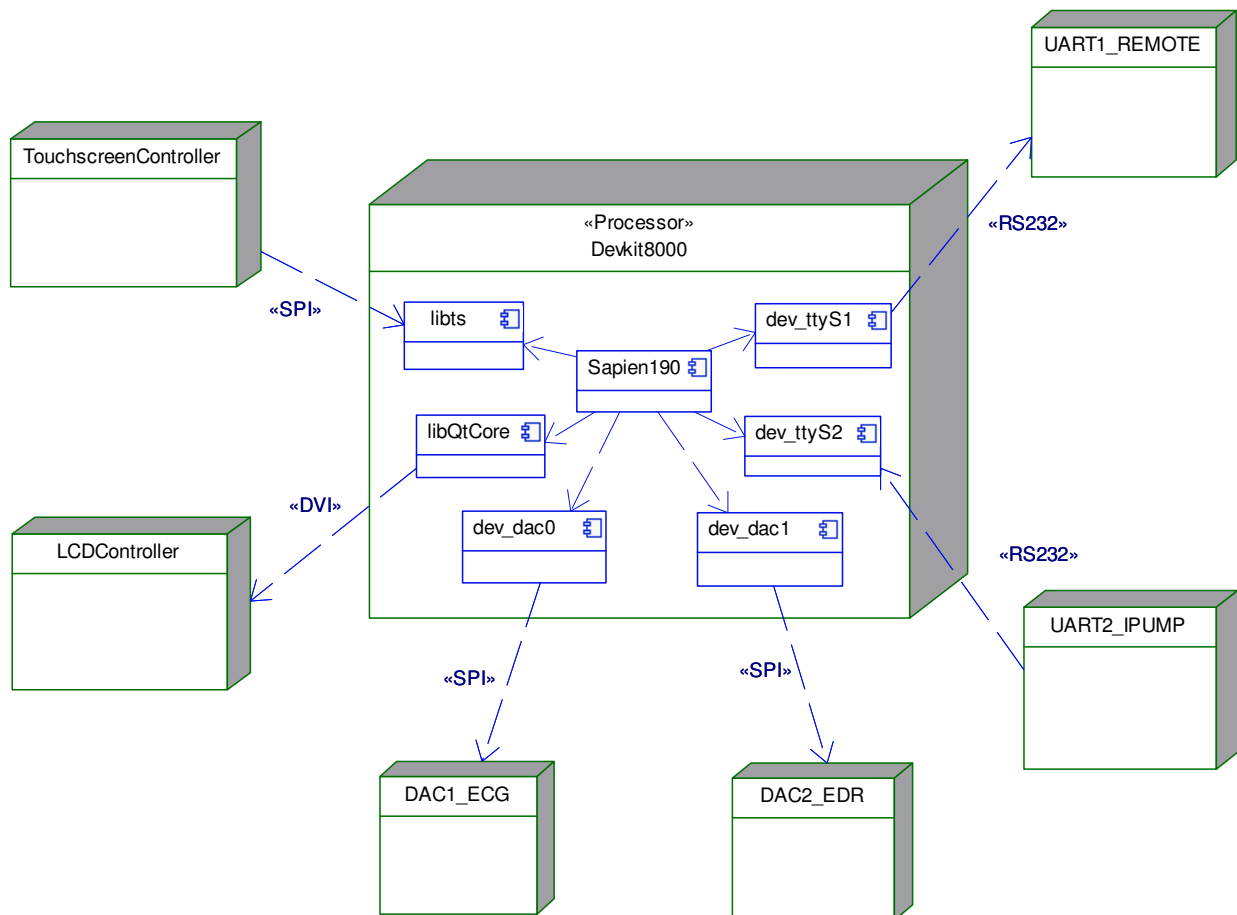


Figure 11 Essential HW nodes used from DevKit8000 used for Sapien190

2.3.2. Logical view - discrete package

The figure below illustrates how the command and state pattern is used to implement the state machine for the UI controller in the discrete package. We have modified the state pattern represented with the SimState class and its sub-states. We have implemented a bidirectional association between the context (UI Controller) and state (SimState) to remove passing the reference to the context every time a command is executed and the state is changed. Instead the reference to the context is passed to the state when it is created. The approach has been implemented for the command pattern. Every time a command is created an association to the state is established and thereby saves passing a parameter to the execute method in the command. Since this part of the design is using a number of Qt classes and primitives like slots and signals we have only used Rhapsody as a drawing tool for this part of the design see complete UML class diagram for design of the command and state pattern.

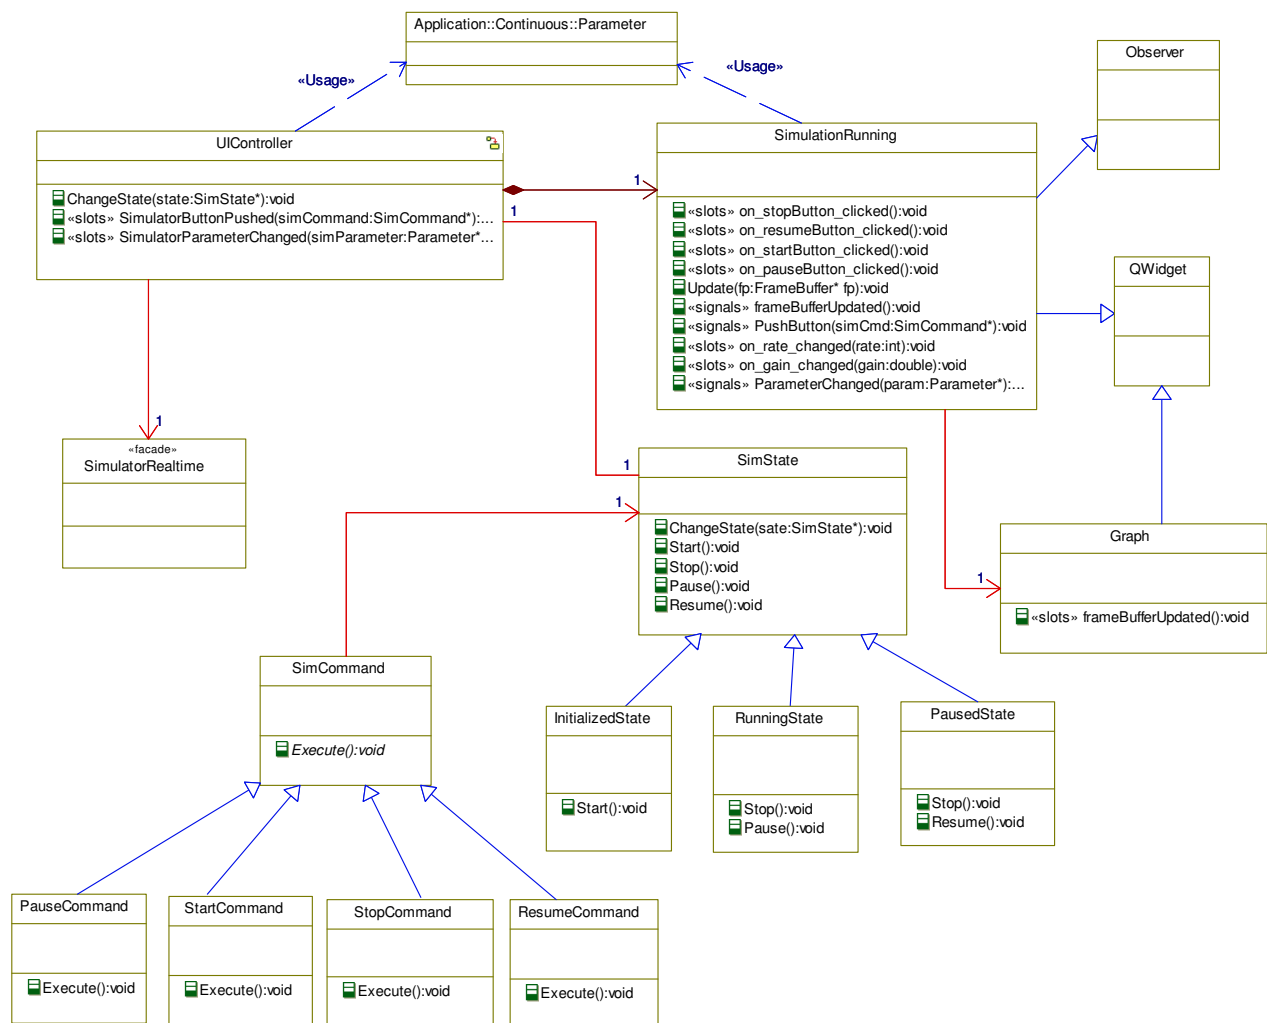


Figure 12 Command, State and Observer pattern used to design SapienApplication (Controller)

The discreet package provides complete control of the continuous system and processing and displaying the output coming from the discrete system.

Our continuous system has three states: Running Paused and Stopped. In managing this, we chose to implement state pattern, SimState class seen in Figure 12. The State pattern is a nice way to keep track of conditions and the opportunities available in the states. For instance, it should not be possible to start the simulation when it is paused, here you must instead use resume.

The Qt framework offers a form of call-back when working with the UI framework. This is used together with the command pattern to manage our state machine. Command patterns advantage is that you can define a number of things to be done when the function execute is called. Say that we want to stop the simulator so we set up a stop command object, which only purpose is to execute and run the function stop on our state machine. This means that we have stream lined the way commands from the UI will be executed in our continuous system.

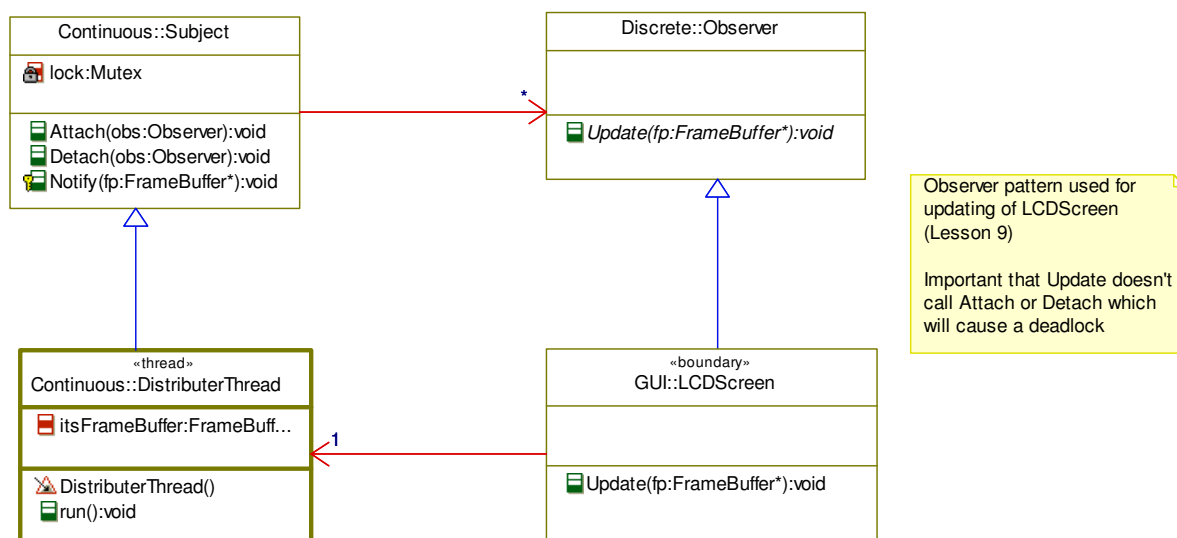


Figure 13 Observer pattern used to notify GUI with new frame buffer

To carry information from the continuous system have we used the observer pattern, as shown in Figure 13 This pattern make it easy for sub-systems who want something specific data to be made aware of when new data arrives. In our system it is used so our UI update the graph when the continuous system has generated any new data.

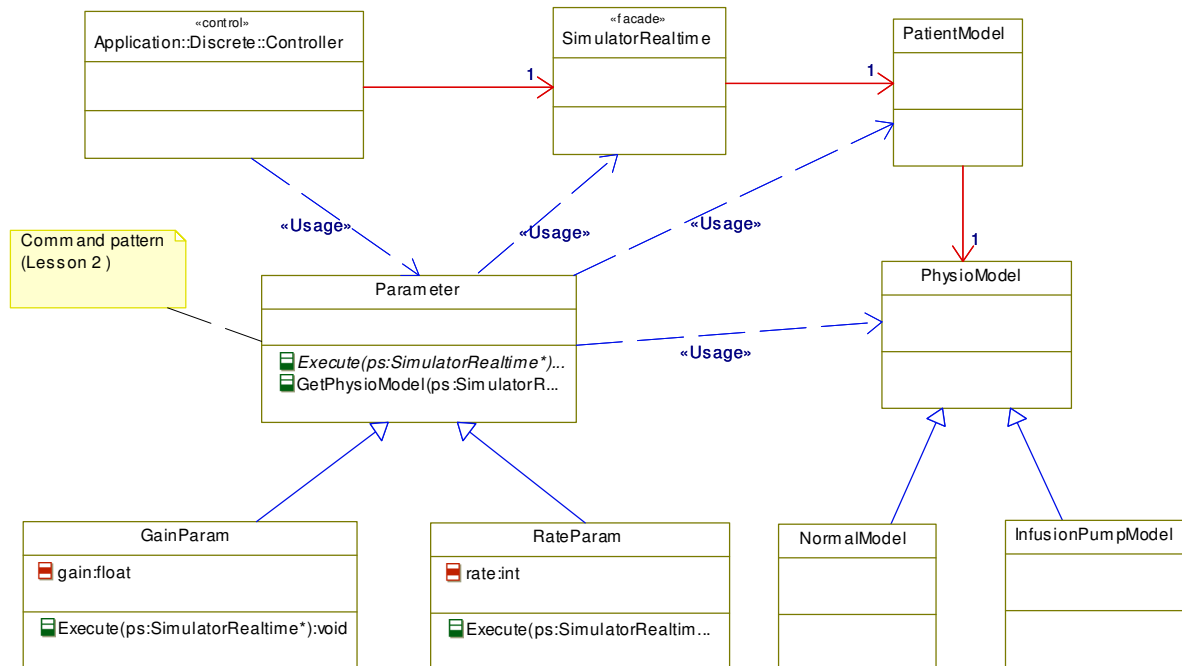


Figure 14 Command pattern used to set parameters in real-time simulator

Like we have used command pattern to control our state machine, we used command pattern to control the way we change our parameters, this is done also to streamline the way this is done.

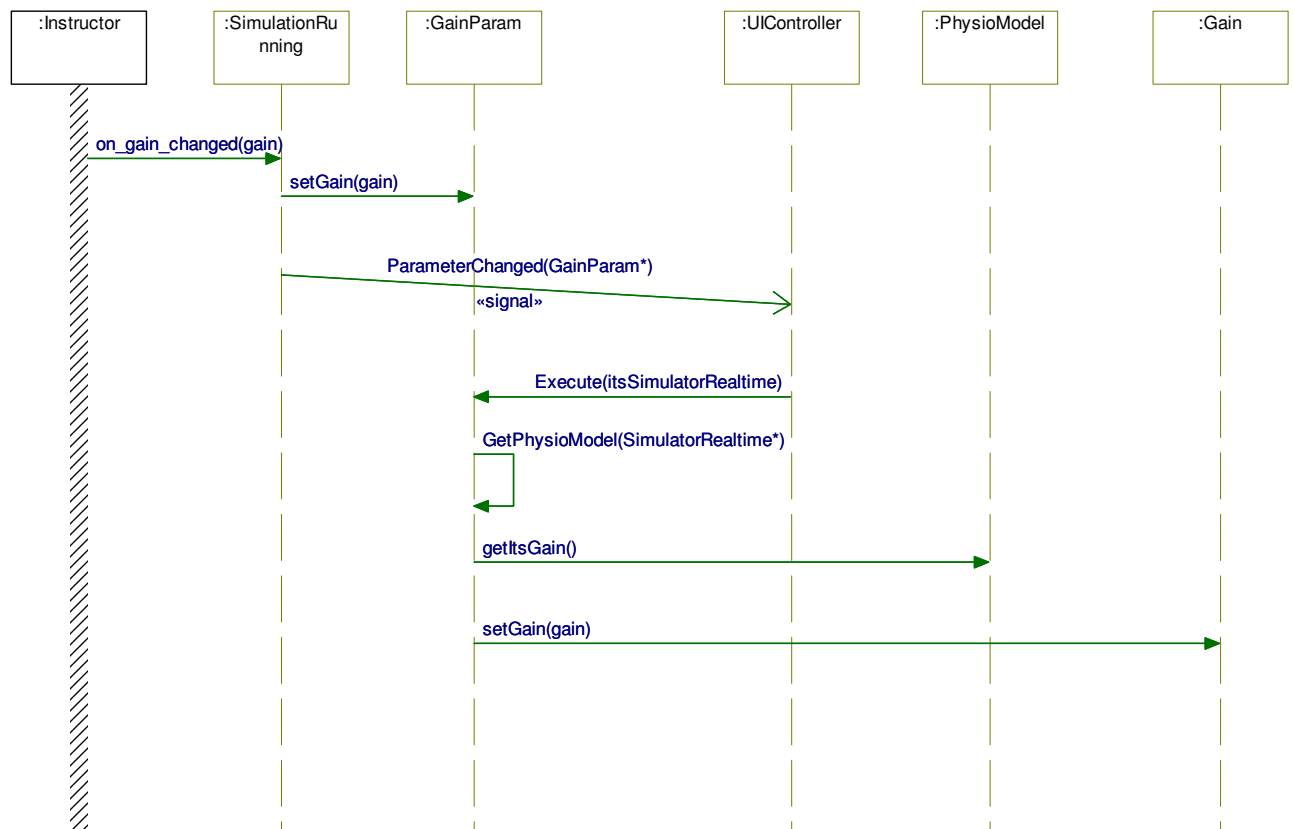


Figure 15 Scenario for instructor adjusting gain parameter

In Figure 15 we see a sequence of what happens when we change one parameter in the UI and how it will be changed down in our continuous system. The instructor changes the parameter of UI. The gain spin slider is change and the `on_gain_changed` method is invoked setting the gain "command" parameter. The gain is set and a callback is emitted asynchronously to the UIController by use of the signals and slots provided by Qt. Actually Qt are in this situation using a combination of an observer and message queue pattern.

2.3.3. Logical view - continuous package

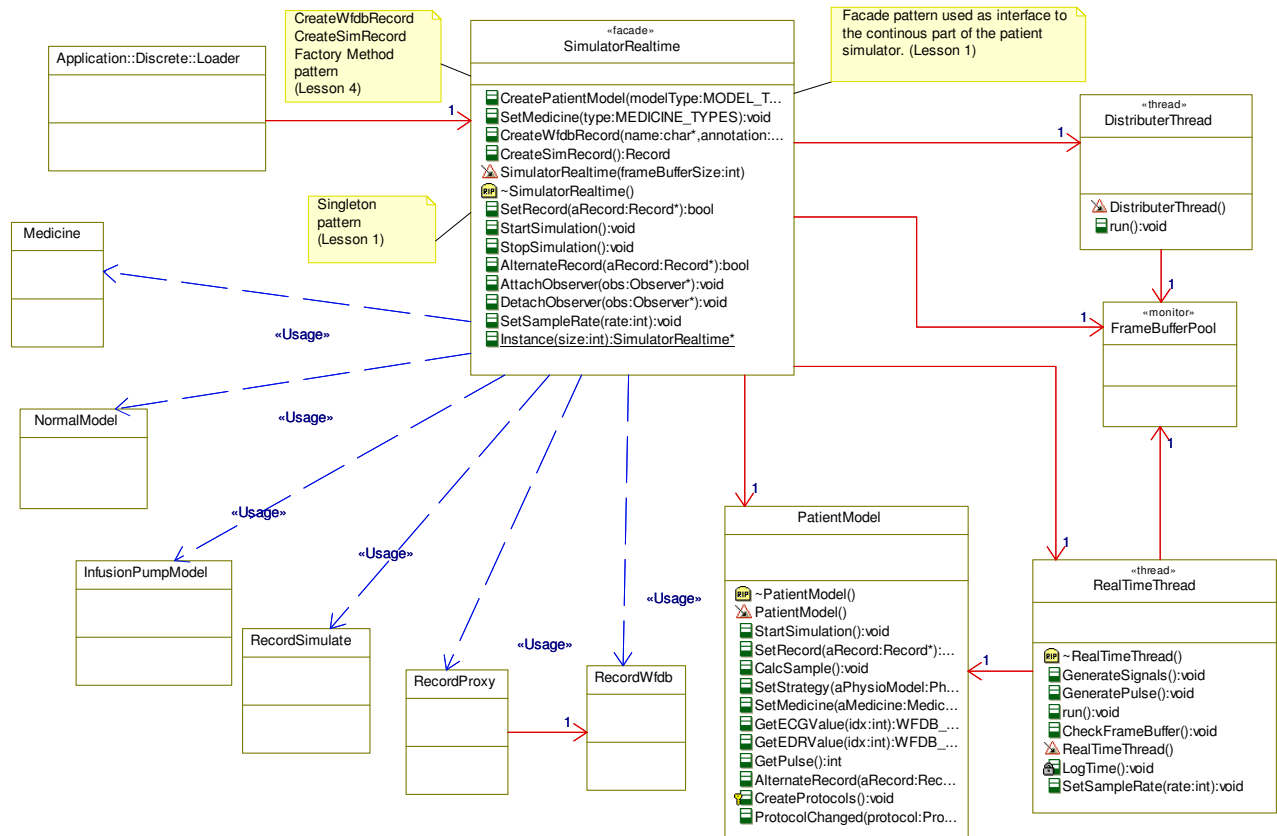


Figure 16 Façade Pattern used for interface to the Continuous Package

The continuous package handles all real-time related work and thus contains a lot of functionality. We could have let the discrete package communicate directly with objects it needs access to, for example use the methods in recordWfdb to set the current Wfdb record in use. This would however create a lot of bindings and go against the OCP principle¹⁴. Instead we have chosen to apply the façade pattern (Gamma, 1994 p.185) to the whole continuous package. The façade pattern provides a single class for the outside to interface to. The complexity behind the façade is hidden to the outside. In our case the façade becomes the entrance point for control and configuration of the discrete system. Since this is the case, the façade should also be responsible for creating the involved objects, as it holds their initialization data. This is described in the "Creator Pattern" known from GRASP (Larman, 2008 p. 282).

¹⁴ Martin, Robert C. "The Open-Closed Principle" 1996
<http://www.objectmentor.com/resources/articles/ocp.pdf>

One of the key problems defined in the introduction, is how to create a design that is easy to extend and maintain. For a system like the patient simulator, it is obvious to keep the system open to new data input sources. This has been accomplished in two ways: By using the proxy pattern to allow adaptation to new data sources and by using the iterator pattern to move sample sets from the record to the patient model.

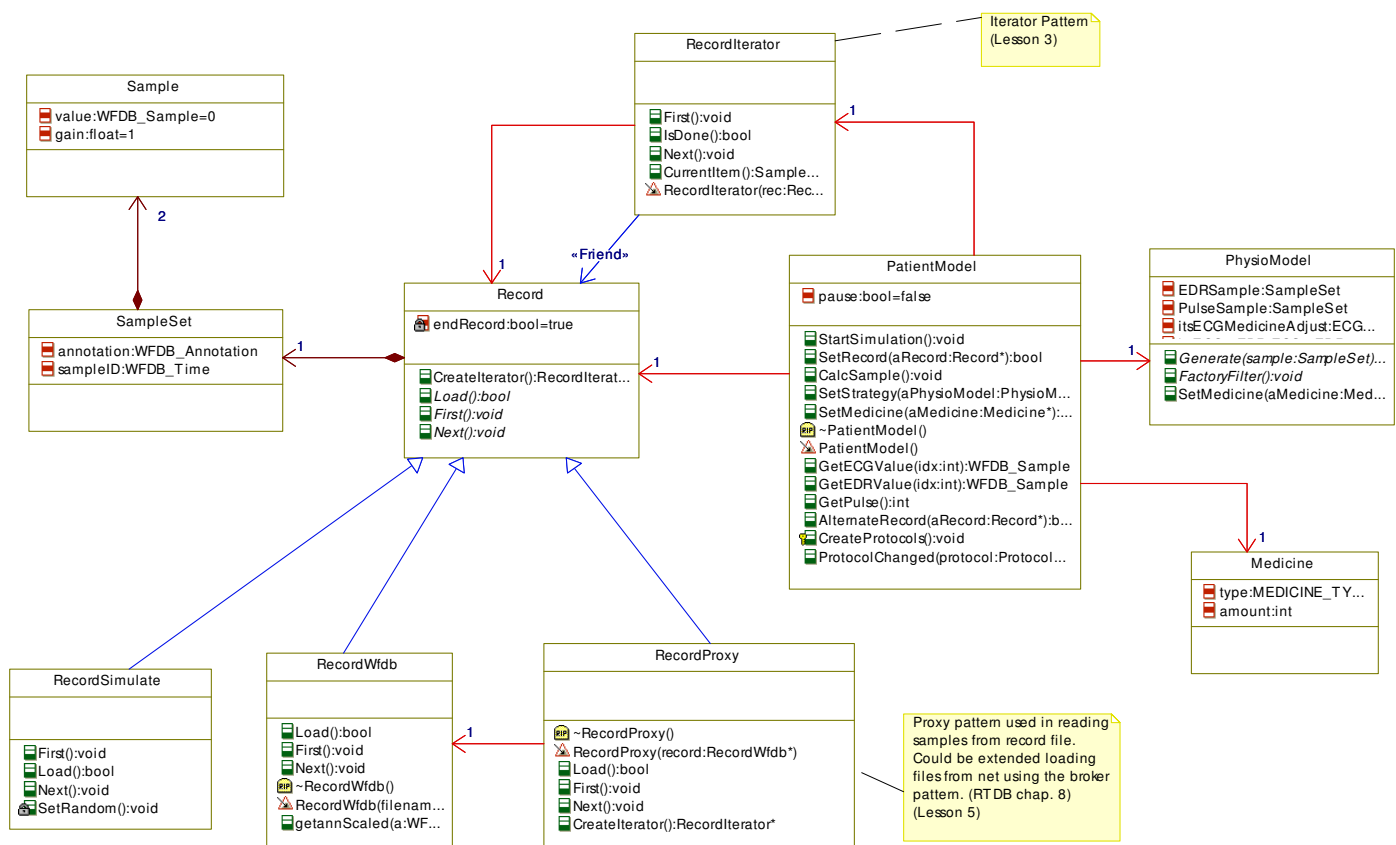


Figure 17 Proxy and Iterator Pattern used to access Records from the PatientModel

The proxy simply acts as a placeholder for a future data source. It provides indirection of *Record* class' methods and complies with the Liskov Substitution Principle¹⁵.

The Iterator Pattern was chosen because it decouples the *Record* and the *PatientModel* and supports variations in the objects passed between the two classes. The iterator pattern (Gamma, 1994 p. 257) uses the "Pure Fabrication" pattern (Larman, 2008 p. 421) by creating an "artificial" class (the Iterator) to take responsibility of moving data, instead of leaving it to *Record* or *PatientModel* and thus increasing the binding between the two. Our implementation uses an External Iterator implementation, as the *PatientModel* sinks data directly from the *Record* source.

¹⁵ Liskov, Barbara. "Data Abstraction and hierarchy" 1987

Another important place to be open for extensions is the interface to external signals that may have an impact on the physiological model. The current version of the software does not support an actual infusion pump, that would be able to inject medicine and thus change parameters in the physiological model. The design is however prepared for this. The *PhysioModel* currently has two strategies, one with the infusion pump included and another without. If additional sources are requested, then the design can be extended, but it needs not to be modified.

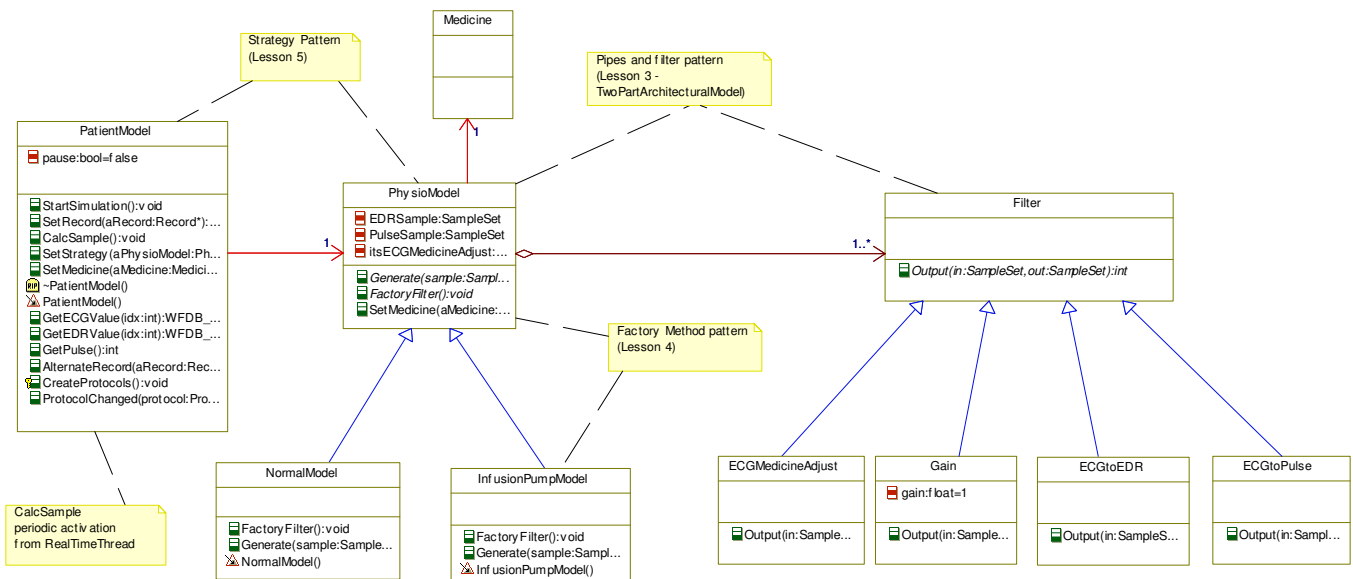


Figure 18 Strategy, Filter and Pipes Pattern used for PhysioModel

The factory method pattern is used to instantiate the actual filter implementation. Each strategy implements its own version.

Yet another place to be open for extensions but closed for modifications is the *PhysioModel* itself. The model consists of a number of digital filters that is combined to create a physiological model. The implementation is based on the pipes and filters pattern (Shaw, 1996). Using this pattern, each filter appears as a 2-port black box, allowing us to combine any sequence of filters to create a desired result. New filter types can be added, as long as their interface conforms to the two-port interface. The filters are combined in the factory methods mentioned previously. Figure 19 shows an example of the *PhysioModel* in action.

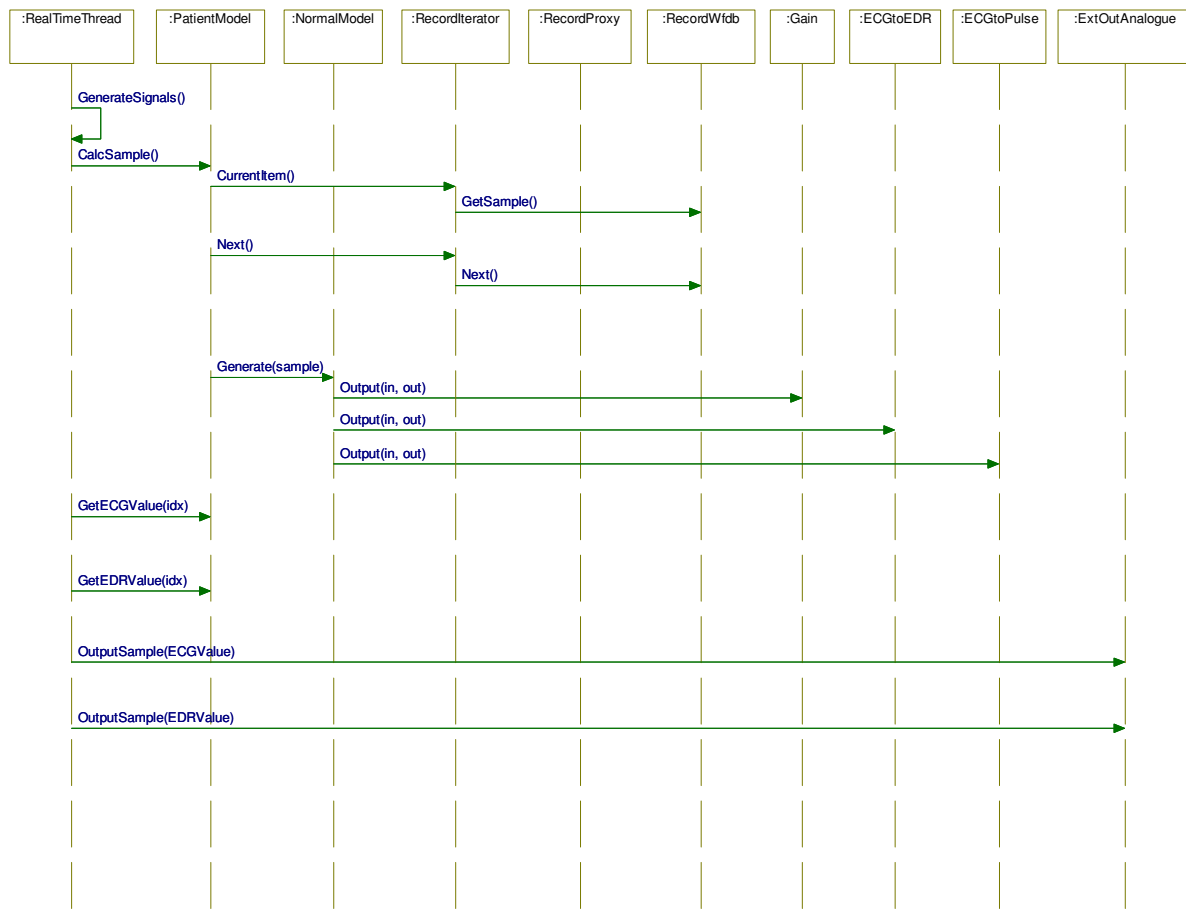


Figure 19 Sequence Diagram: GenerateSignals using filter and pipes

Figure 19 shows the iterator come into action to get a record data set. The Filters and Pipes pattern is applied to the data set. The result of the data processing is output to the abstracted hardware interface.

2.3.4. Logical view - communication package

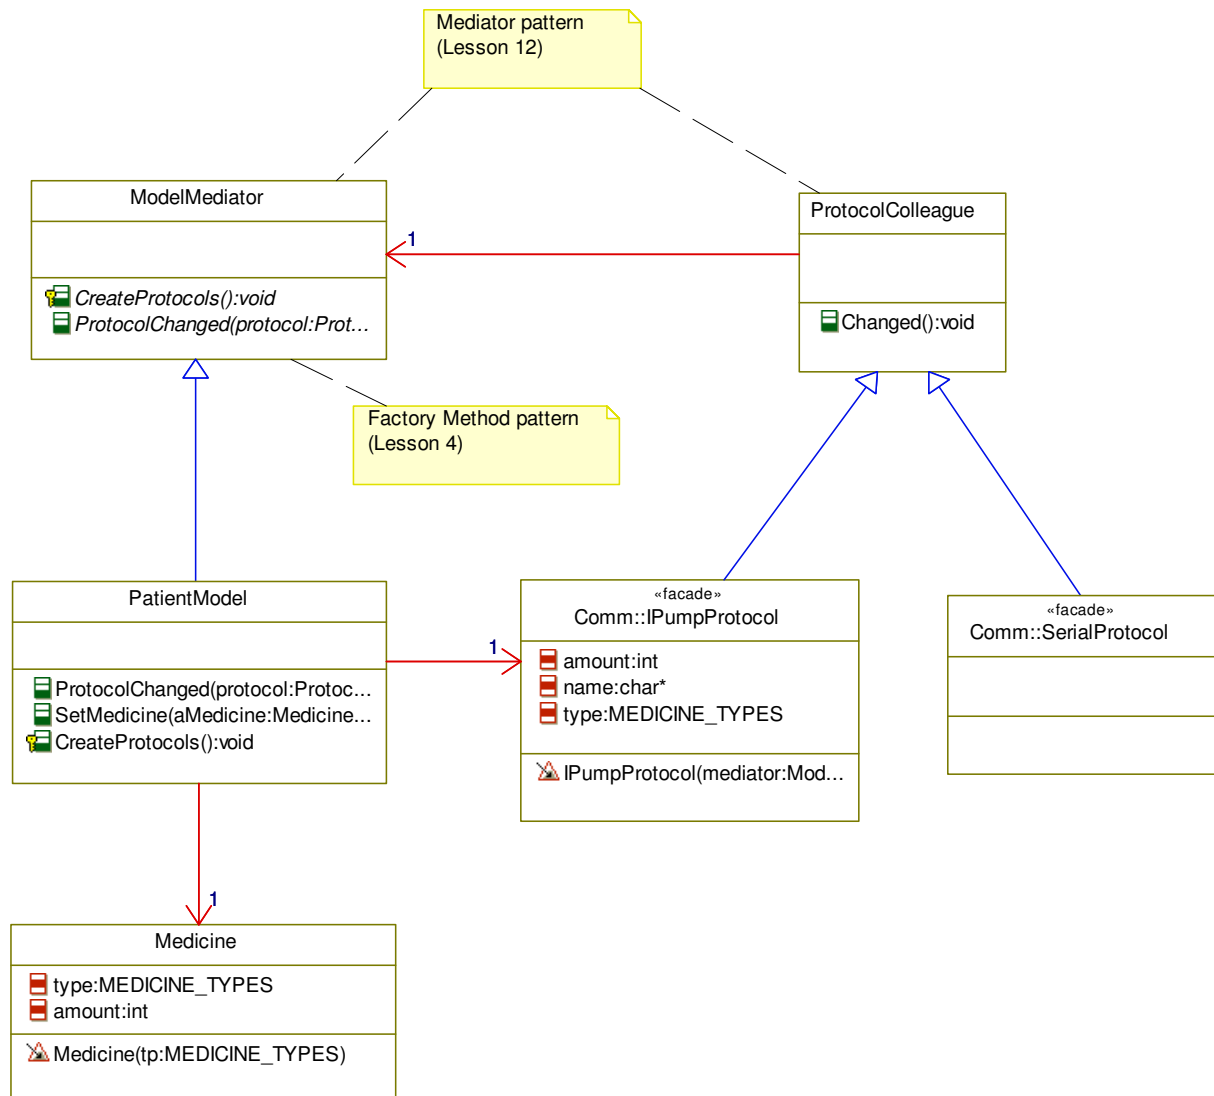


Figure 20 Mediator pattern used to update PatientModel with external input from protocols

To control external input to our patient model, we designed a mediator pattern. This has not been implemented yet, just design. The idea of the Mediator pattern is that it encapsulates how objects interact inside one object. This is to promote low coupling by allowing the objects through the same object. In our case we have chosen to use to prepare for future external inputs. It therefore ensure us that if we add more inputs, it can easily be transported down to our patient model without having to make major changes in existing code.

2.3.5. Process view

In this chapter we will briefly describe the process view of the Sapien 190 simulator. The following threads are running in patient simulator:

Controller:

The controller will be part of the discrete package that contains the state machine to control the continuous part of the two layered application model. Synchronization is needed between the controller and the continuous part of the system.

RealTimeThread:

This thread is the essential thread of the systems that is periodic with the sample rate and is responsible for generating signals and outputs them to the environment (Push). Alternatively this thread could have been implemented creating a buffer of samples and sending it to the DAC driver. The DAC driver should then be responsible for sending the buffer at a given sample rate (Pull). This approach would have been more efficient and accurate, since timing in kernel space of Linux would be more predictable in calculation of WCE. It also introduces less overhead in copy of data from user to kernel space.

DistributerThread:

This thread is used to collect a frame buffer of samples that is updated to the observers which in this case will be the controller class part of the "Qt-GUI thread". The DistributerThread is periodic with the sample rate times the number of samples in the frame buffer.

IpumpThread:

This thread is used to handle PDU messages received from the IPUMP and updating the medicine volume. PDU messages are received with a rate of 1 sec.

The controller class from the discrete part of the system will interact with the DistributerThread where the observer pattern has been used to update the GUI with frames of samples. A mutex has been added to this observer pattern to ensure synchronization between the controller thread and the DistributerThread. This synchronization will ensure that Attach is not called the same time as Notify, meaning that new observes are not allowed to be added or deleted at the same time we are updating the observes. The RealTimeThread generates samples and collects a number of samples before a new frame buffer are sent to the DistributerThread by using a mailbox (SendMail and GetMail). The FrameBufferPool is implemented as monitor to ensure synchronization between the RealTimeThread and DistributerThreads when they request the pool of free frame buffers.

Synchronization between the IpumpThread and RealTimeThread is not yet finalized since this part is still in the design phase. It could be done by adding a monitor or mutex to the mediator pattern to ensure that the medicine class is updated and access exclusively between the two threads.

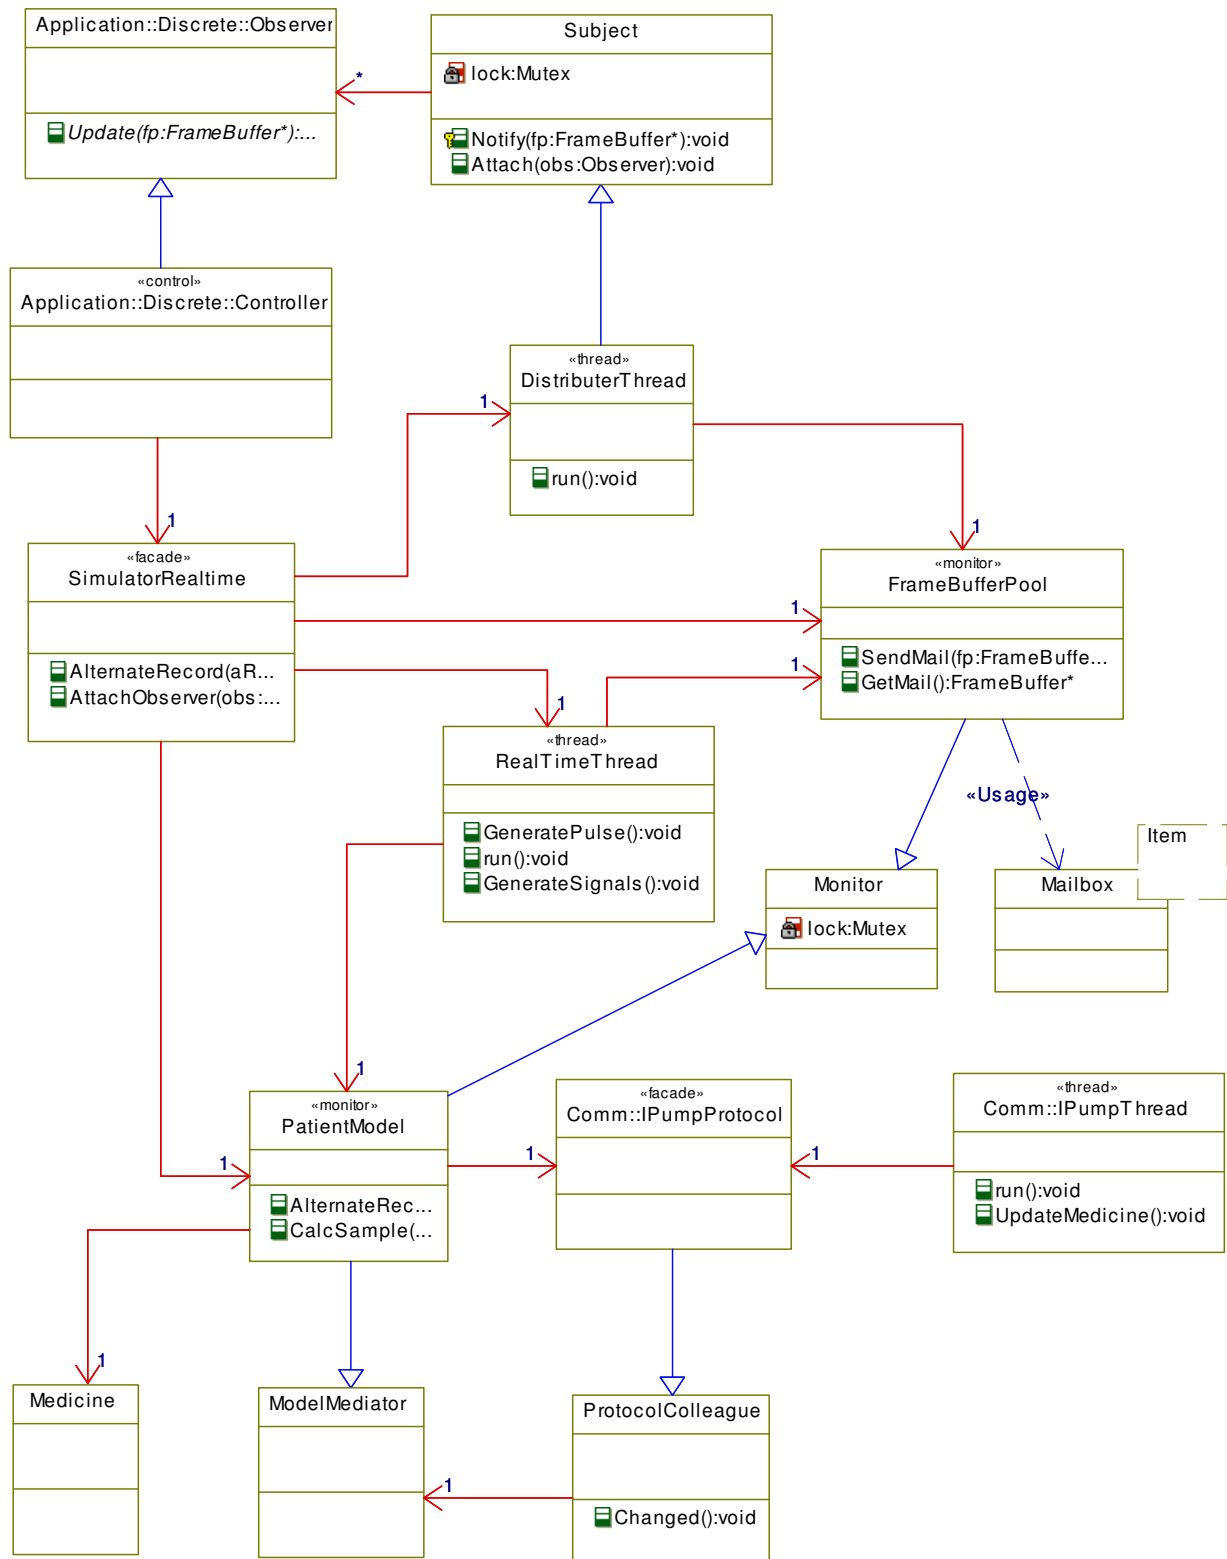


Figure 21 Overview of threads for the Sapien 190 design

The essential threads for Sapien 190 are the DistributerThread and RealTimeThread. These threads exchanges frame buffers with samples by use of a Mailbox implemented according to the Message Queuing Pattern. This pattern uses asynchronous communication between the two threads. The FrameBufferPool is implemented according to the Pool Allocation Pattern. This approach saves allocation of a new frame buffer from the heap every time the RealTimeThread will be filling the next buffer with samples. There is only allocated 2 frame buffers in the pool since the DistributerThread needs to distribute the newest samples to the waveform graph. In case it cannot follow the speed of updating the graph if the sampling rate is too high samples will automatically be skipped. We could also have used the Static¹⁶ Allocation Pattern to achieve the same effect, but the GenericPool class is a more generic solution that can be used in other parts of the design for future extensions. The monitor is used by the FrameBufferPool to synchronize allocation and release of FrameBuffer's to the GenericPool.

¹⁶ Slide 3 Lesson 6 – MemoryPatterns.pdf

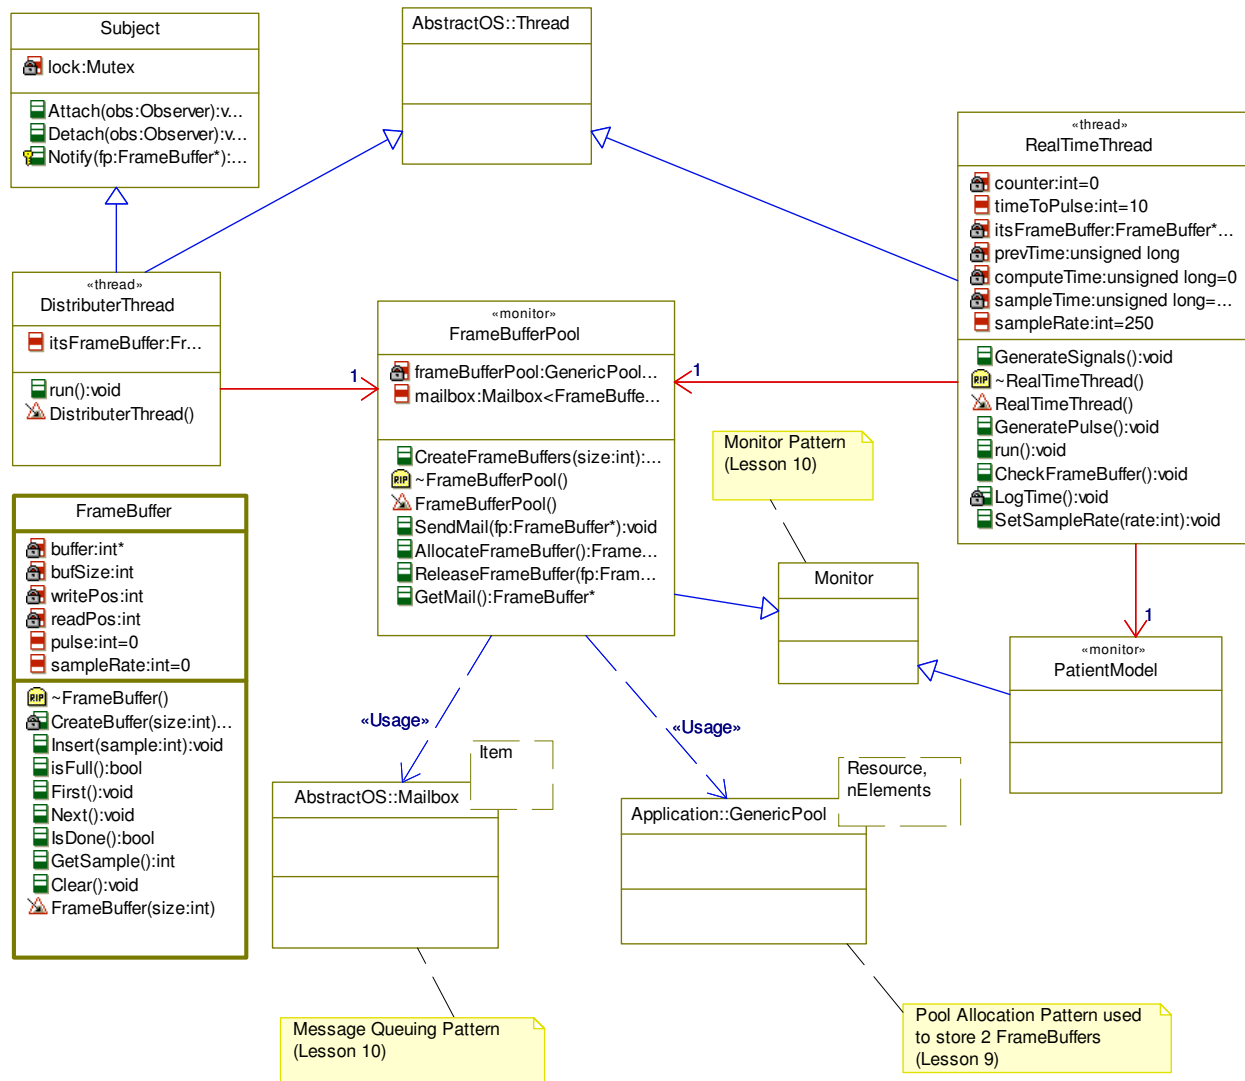


Figure 22 Process view for Distributer and RealTime threads and mechanism for synchronization

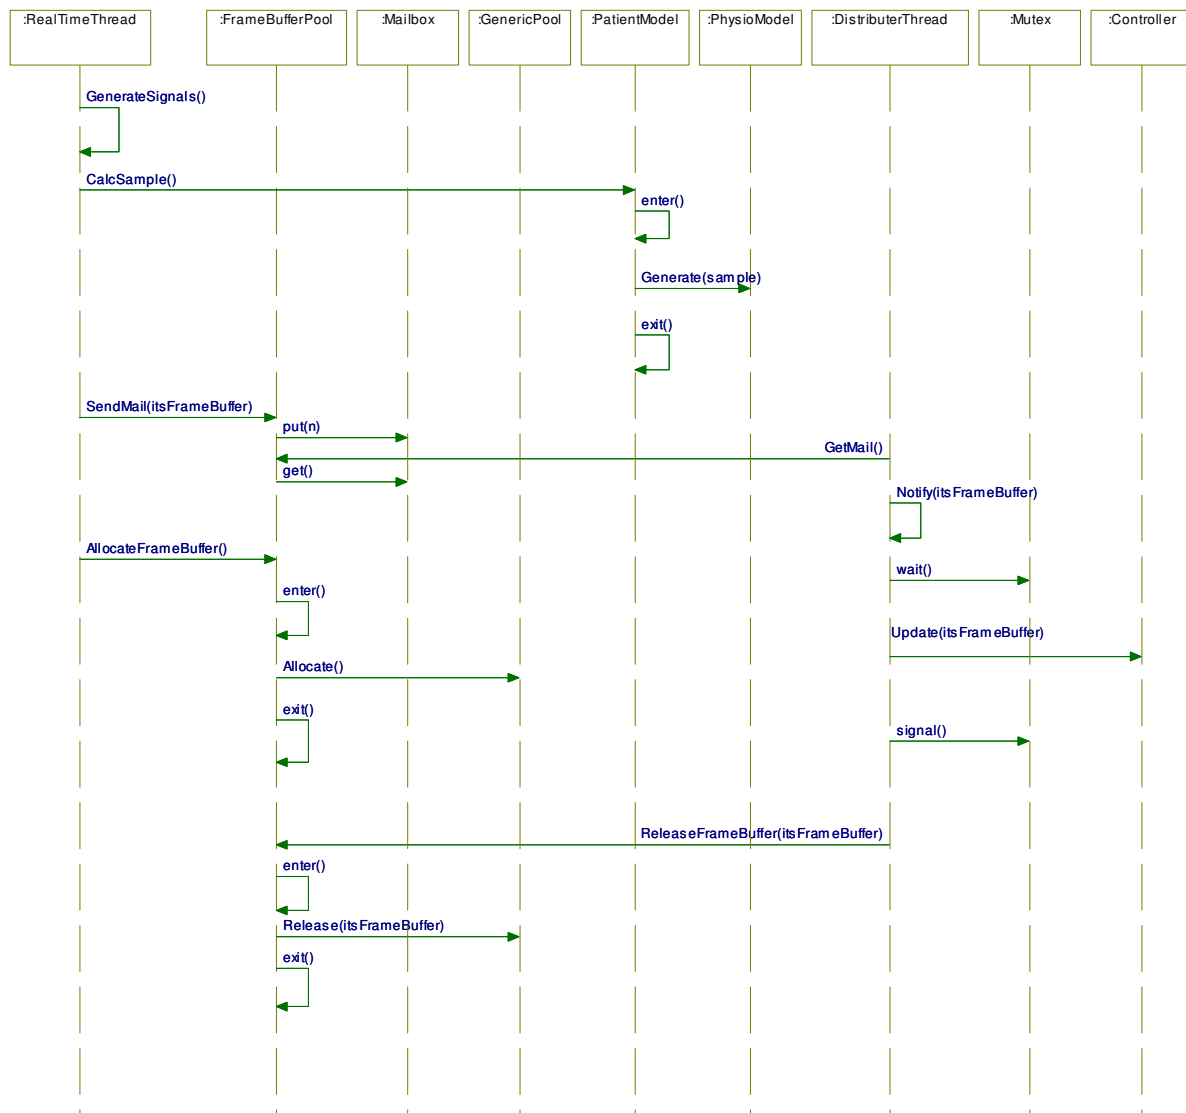


Figure 23 Synchronization between threads

The operating system is encapsulated in the classes displayed below. These classes are implemented in two versions. The first is to be used for simulation and test in Rhapsody. The second is implemented as an abstraction of the POSIX thread API used on Linux.

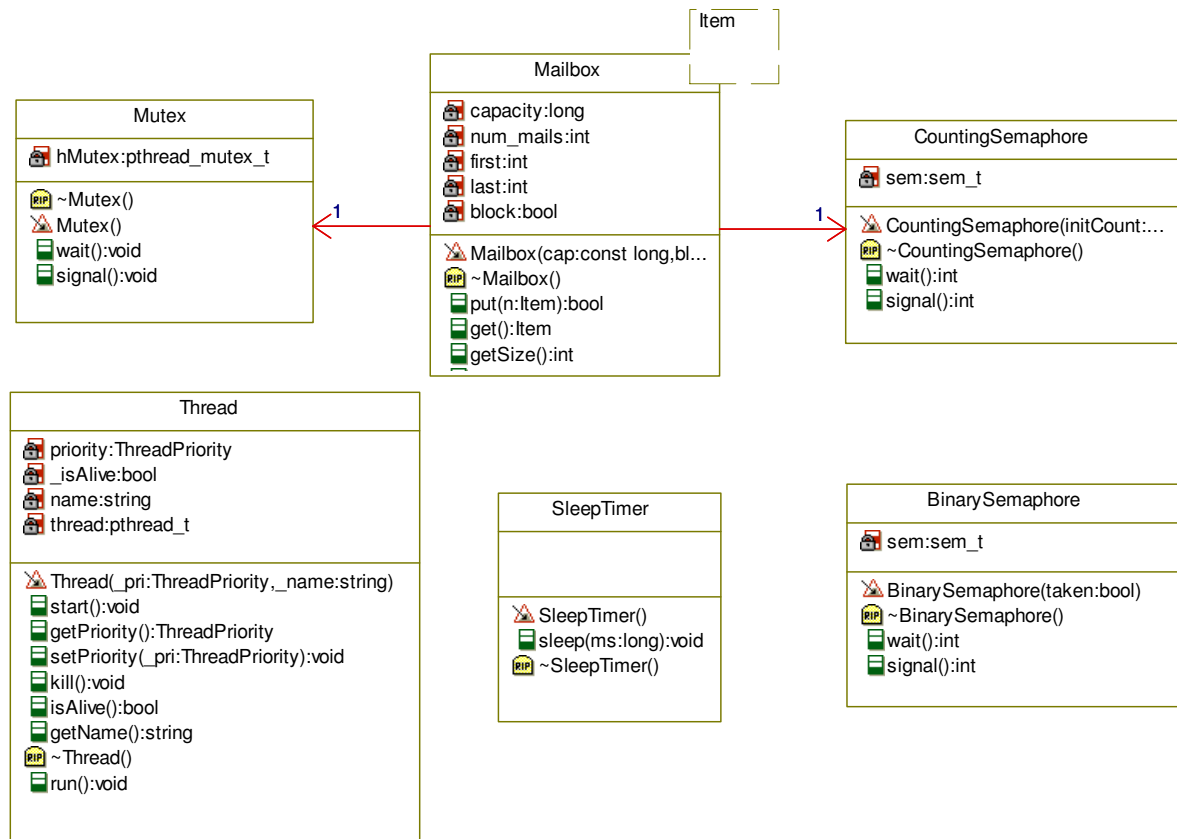


Figure 24 Abstract OS (Linux)

2.3.6. Rate Monotonic Analysis

Based on the event analysis made in chapter 2.2 we have included a Rate Monotonic Analysis of the scheduling of threads. The WCE times and blocking delays is based on measurements and estimates. We have measured the time it takes to generate new ECG and EDR samples WCE (250) and the time it takes to update the waveform graph to only ~4000 us instead of the 100 ms (B_i for DistributerThread) used in the RMA below. All times below is in us.

#	Event Id	Arrival Period (T_i)	Action	Thread
1	Sample	4000	Run (GenerateSignals)	RealTimeThread
2	Pulse	40000	Run (GeneratePulse)	RealTimeThread
3	PDU	1000000	Run (UpdateMedicine)	IPumpThread
4	FrameBuffer	200000	Run (Notify)	DistributerThread

WCE Time (C_i)	Priority	Blocking Delays (B_i)	Blocking term ($B_{ti} = B_i/T_i$)	Deadline (D_i)
250	Very High	40	0.01	4000
50	Very High	0	0	4000
200	High	40	0.00004	10000
200	Medium	100000	0.5	200000

In the below calculation we can see that with a sample rate of 250 Hz we will be able to schedule the threads using a frame buffer of 50 samples. We can see that the calculated total utilization (UB_{total}) is less than the utilization bound (U_{bound}).

Rate Monotonic Analysis with Task Blocking

(U = Utilization, UB = Utilization and Blocking, B = Blocking)

U_{total}	$\sum (C_i/T_i)$	0.06
U_{bound}	$n(2^{1/n} - 1)$	0.76
B_{total}	$\max(B_{ti})$	0.50
UB_{total}	$U_{total} + B_{total}$	0.56

$U_{bound} > UB_{total}$

In the next part we have added calculation of the utilization bound for the FrameBuffer event, since this is the part that is most complicated and critical for the scheduling updating the waveform graph. More details can be found in chapter 6.5 of the architecture document reference [4].

Utilization bound for FrameBuffer (e4) valid as long $N_f > N_p$

<u>Step 1: Identify H</u>		<u>Step 2: Calculate f</u>
- Higher priority events - e1, e2, e3		
$H1 = e3$	$T_i \geq Di(4)$	$f4 = \sum(C_j/T_j) H_n + 1/T_i(4) (C_i(4) + B_i(4) + \sum(C_k) H1$
$H_n = e1, e2$	$T_i < Di(4)$	0.57

<u>Step 3: Utilization bound</u>	
$u(n, di) = n((2*di)^{1/n} - 1) + 1 - di$	0.83 (0.5 < di <= 1)
$di = Di(4) / Ti(4)$	1.00 (di < 0.5)

<u>Step 4: Compare effective calculated utilization with bound</u>	
Calculate $f < Utilization\ bound$ ($f4 < di$ or $u(n, di)$)	TRUE

2.4. Translation and testing

Rhapsody (version 7.5) has been used to create an UML model for the Sapien 190 patient simulator that is used for test of the simulator model before automatically generated C++ source code to be compiled on Linux and target. The Model Driven Architecture (MDA¹⁷) approach has been used where we have created a platform independent model (PIM – on Windows) that later is translated to the platform specific model (PSM – on Linux) target the Linux platform by using the code generation from Rhapsody. This approach has fully been used for the continuous package of our design.

We have created test scenarios by use of Rhapsody state diagrams. In Rhapsody it is possible to set breakpoints in the animated state diagrams and perform an inspection of the state variables of the model (Instances of classes and attributes) by using the high level visual abstraction in Rhapsody. During execution of the model animated sequence diagrams has been created to verify the design

¹⁷ Slides 31 – 34 from Lesson 8 – L3_ROPES_Process.pdf

model see example in Figure 26. This approach has reduced the development time in removing a lot of the manual C++ coding and debugging of the model. The translation from UML to C++ code is fully automated. The contents of methods in Rhapsody still need to be coded manual in the UML model.

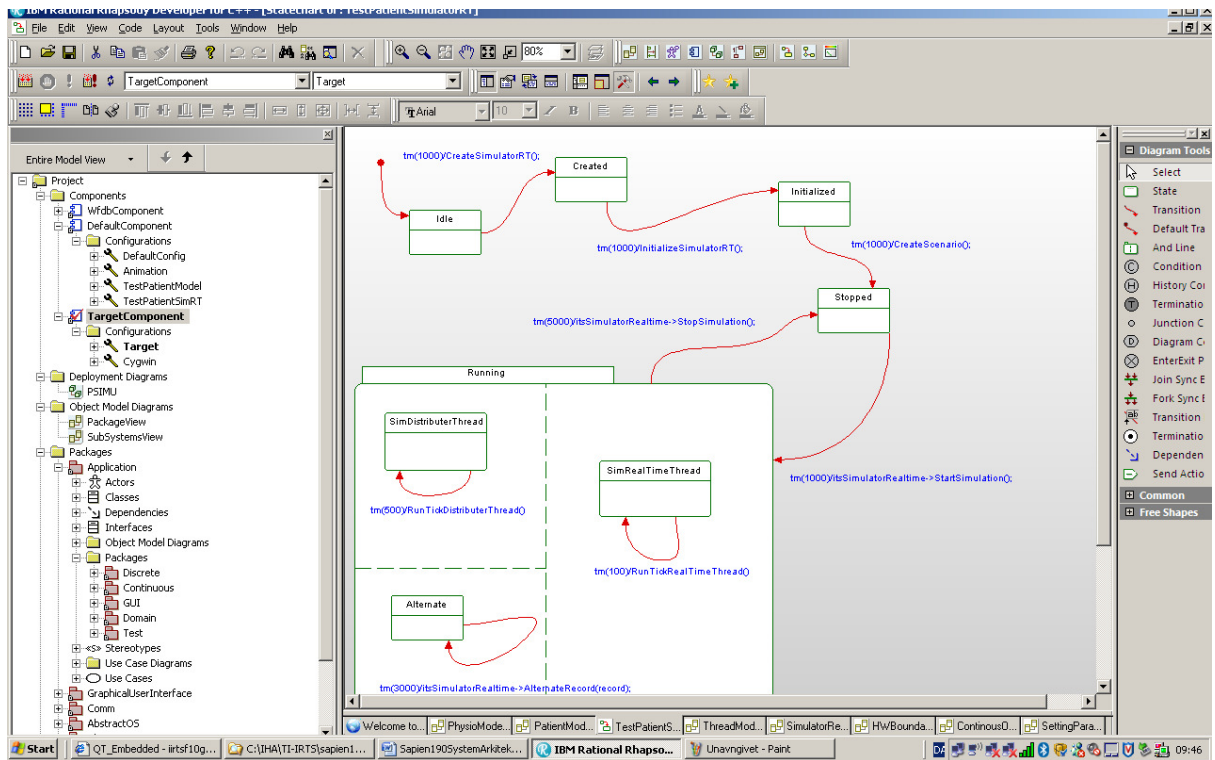


Figure 25 Rhapsody UML model used for simulation and testing

Figure 26 Animated sequence diagram for testing in Rhapsody

host or target. More details on compilation and installation are to be found in chapter 12 and 13 of the architecture document see reference [4].

2.5. Development tools

2.5.1. QT Creator

We have chosen to use the Qt Creator development environment to support development of graphical user interfaces target different embedded platforms including Linux.

2.5.2. Rhapsody

For implementation of the architecture itself, we have used Rhapsody. When a system is first created in Rhapsody, it is very easy to develop further on it, and test new functionalities.

2.5.3. Eclipse

To start off with, we used eclipse but when we had to make the user interface we were forced to switch to Qt Creator. However, we later found out that it is possible to get a plug-in for eclipse that allows us to develop in the eclipse environment by using the Qt Framework.

2.5.4. Subversion and Google Code

Throughout the development we have used Subversion (SVN) for version control. We have chosen to use Google Code as SVN server. It is possible to get a svn client for Linux and windows so it is also a good tool to share code across the two platforms

2.6. Results

The final prototype of the product is running on target and contains the basic and essential functionality for the real-time part of the product covering functionality based on the use cases #1, #2 and #3.

The prototype contains a graphical user interface in where the simulator can be started and stopped. When the simulation is running it is possible to alternating playing different records. Parameters can be adjusted during the simulation by setting the gain to zero simulating a heart stop. The rate of samples send out to the DAC can be adjusted to measure performance limits. The signal waveform graph of replaying the ECG records will be simultaneously updated on the LCD screen at the rate of each 50 samples. The pulse is computed and output is displayed on the serial console.

The design has been tested in Rhapsody before testing on the Linux host platform and final testing on target. Details about the size, performance and quality measurements are to be found in chapters 10 and 11 of the architecture document 4.

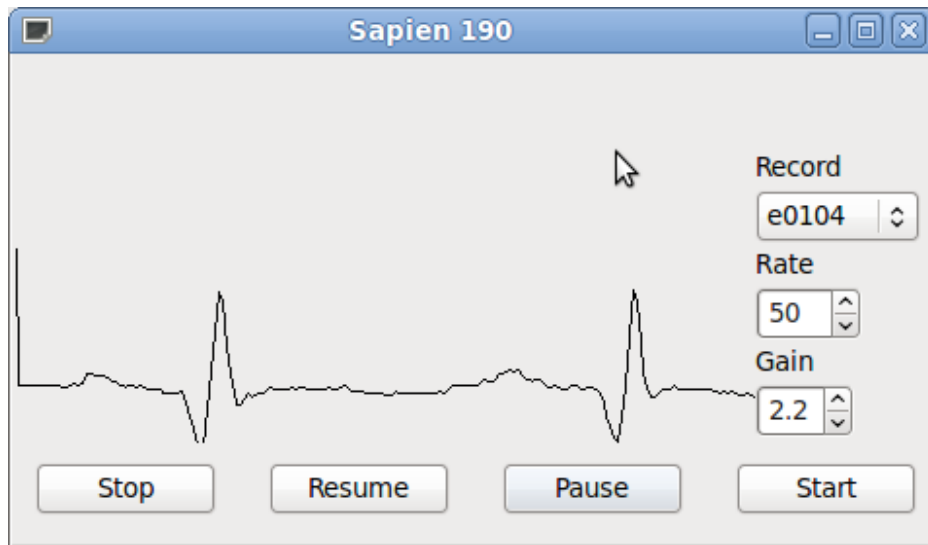


Figure 27 Sapien 190 Final prototype

The figure above shows the GUI of the patient Simulator updating the ECG waveform graph. Below is listed the serial console output that generates the calculated pulse.

```
Pulse 52  
Pulse 55  
Pulse 60  
Pulse 58  
Pulse 62  
Pulse 62  
Pulse 61
```

The figure below shows the ECG signal from the analogue output from the patient Simulator, though it is rather noisy. The noise issues are electrical related and not within the scope of this project.

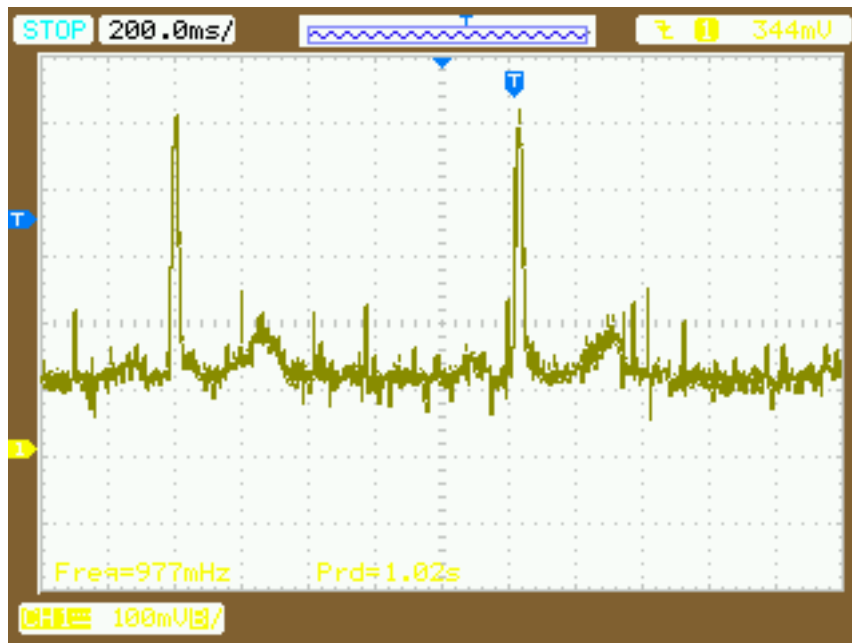


Figure 28 DAC Output ECG

It is important that refreshing the display does not take too long time. If the update rate becomes too slow the display will appear to be flickering and will become uncomfortable to look at. These measurements will also be used in the RMA analyses to calculate if thread scheduling would be possible. The update and compute times are measured by adding code to log time (us) in Linux. Measurements are calculated for updating the waveform graph and the time the RealTimeThread takes to compute and generate output samples (ECG, EDR and pulse).




Figure 29 Graph Update Time

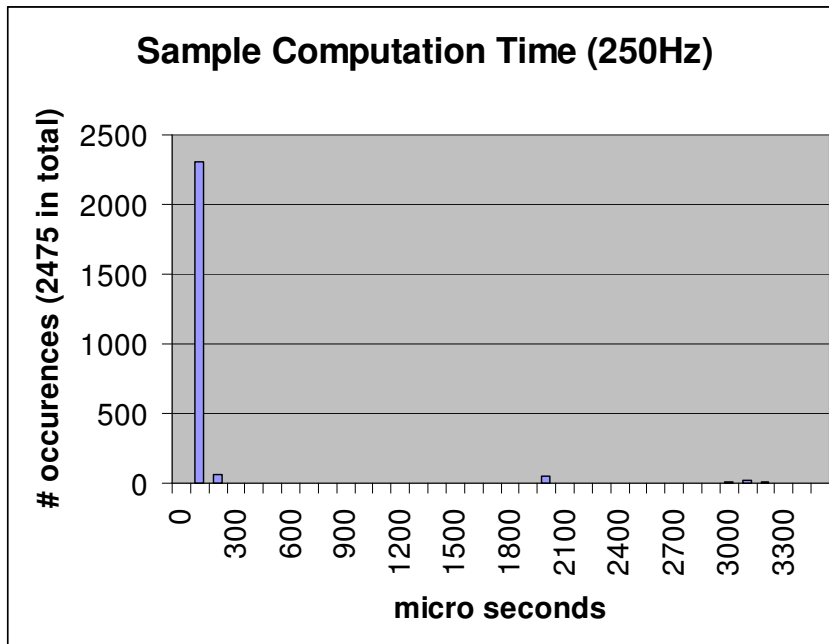


Figure 30 Sample Computation Time (250Hz)

The required flash disc size by the application and libraries is indicated below:

Application:	
Sapien190	0.15 MB
Libraries:	
libQtCore.so.4	3.10 MB
libQtGui.so.4	11.00 MB
libQtNetwork.so.4	0.89 MB
libwfdb.so.10	0.16 MB
libstdc++.so.6	1.18 MB
Libts-0.0.so.0.1.1	0.01 MB
Total:	
Lib+Application	<u>16.50 MB</u>

Additional the WFDB record files should be counted. A 30-minute file has a size of approximately 5 MB at 250 Hz sample rate.

We have also measured the memory usage to 36MB equal to 29% of the memory available on the target platform. This rather massive memory usage is caused by the following:

Libraries:	16.45 MB
Application	0.15 MB
WFDB record file	16.00 MB (36-20MB)
Video Frame Buffer	0.16 MB (480 x 320 x 8-bit)
Total	<u>32.76 MB</u>

The CPU load is light only 2% of the CPU. The reason is that the test application only updates the display at a 5Hz rate. This gives a workload of $= (5\text{Hz} * 3150 \text{ us})\% = 1.5 \%$. Currently it is not the compute and generate of new samples that is critical. The current design seems to be efficient enough for the current hardware platform.

2.7. Discussion on achieved results

In the process of making the architecture and design for the patient simulator, we have been inspired by the ROPES and the SCRUM process which has been the fundament for an efficient and structured way of achieving the results.

We have made a prototype of the patient simulator that has incorporated many design patterns that is part of the course TIIRTS where we have covered both traditional GoF patterns combined with real-time patterns to handle concurrency and memory handling. In some areas the product is perhaps over designed like use of the proxy pattern reading records. This design choice is not "the agile way" of working since we are designing for functionality reading records remotely from the Internet not part of the specification. We have achieved a product that contains the solid fundament for the design and architecture of the patient simulator. Many real-time aspects have been covered in the project like design of the continuous packaged including a rate monotonic analysis of the thread scheduling.

The design is optimized for maintainability, correctness and usability as stated in the requirement specification document. The maintainability is realized use of design patterns like it would be easy to add new patient models by use of the strategy and filter and pipes patterns. It would be easy to add new medicine and parameters realized using the command pattern and adding additional views on the signals by use of the observer pattern. The whole design is carefully documented in the product architecture document 4, making it easy to continue working on the product. The software can easily be moved to another platform due to the design abstractions and use of Qt. (HW, OS and GUI abstractions)

We have verified the correctness of the patient simulator prototype by measurement of the signal output from the DAC and the jitter of the signal. The design uses the DAC output as a sink supplied by

samples of the source from the input records. This simple source-sink design protects against data contention and provides the lowest possible jitter as it is driven by the DAC output. Jitter performance could be improved by letting the device driver control the output timing.

The performance and resource requirement of the target platform for the final prototype is quite good only 2% of the CPU performance is used. It turns out to be the memory consumption that is the most critical for the current prototype using 29% of the memory available. Since we could not figure out how to change the priority of threads running on Linux we have discovered that our real-time thread can be skipping its deadline. Since all threads are running with the same priority updating the waveform graph causes that real-time thread to be skipping sample periods very rarely.

We did expect in the start of the project to be able to fully implement all the use cases we originally did specify, but we prioritized to reduce the functionality and use time on getting the use cases implemented that we designed in the second iteration. We also focused on completing a good documentation for the architecture and design. We have tried to make the design with as few limitations as possible and which is easy to extend and port to other platforms. The current hardware platform is in some aspects limited only to handle 2 analogue output ports, but our design could easily be extended to generate and handle many more signals or even more patient simulators running on the same platform.

The choice of using Rhapsody not only for documentation but also for code generation requires education and experience with the modelling tool. Only one of the project participants had the needed qualifications. Later in the project we decided not to use Rhapsody for code generation of the discrete package but only for documentation. In some aspects, using Rhapsody saves development time (Reduces coding errors and speeds up debugging and testing) in other aspects it increases the development time in mastering the tool and being able to do what you want.

2.8. Experience obtained

During this project we have improved our skills in developing of and embedded real-time system based on the OOA/D methodology incorporating design patterns and handling of concurrency. We have managed to use and transform the theory from the TIIRTS course into a first prototype of a product that is running on an embedded Linux target with acceptable performance. We have learned how to write product documentation based on a use case requirement specification and a product architecture document based on the "4+1 view". We have gained experience with part of the ROPES methodology and Scrum process in the way we have structured the project and organized our work and meetings. Part of the ROPES methodology concerns about analysis and design in where we have gained a better practical experience. Especially in the mechanistic phase of the spiral model in where design patterns are applied.

We have gained experience in designing a concurrent system by combining GoF patterns and concurrency patterns. The continuous part of the design is using a number of concurrent patterns like monitor and the message queue pattern to synchronize threads exchanging information in combination with the observer pattern. We have gained insight in how to make event analysis and rate monotonic analysis (RMA) of the concurrent design. We have found that it actually can be used in an iterative development process where WCE measurements on the prototype can be used to verify how close we are to the limits for scheduling of threads in our product. More details can be found in chapter 6 of the architecture document 4.

Development of embedded systems requires that you have a good understanding on the underlying hardware, the development tools and target environment. In this project we have used a large amount of time trying to sort-out cross compiling in Qt and Eclipse and incorporating libraries for compilation to both Linux development host and target. We have learned that it is important to spend time setting up the environment tool properly from the start, instead of trying to make repetitive tinkering.

We have chosen to use Qt as the "Creator" development tool to support developing of graphical embedded applications. It has been a learning experience to see how the Qt framework is using patterns in the way it is extending the C++ classes with slots and signals. The Qt framework is actually using a combination of an observer and message queue pattern we have learn about in the course. More details on this topic can be found in chapter 5.3.2 of the architecture document 4.

We have had problems checking software in- and out of Subversion on the development host platform due to proxy settings. Debugging is very limited and slow in Qt Creator. Debugging information is very

scarce, which sometimes made it unnecessarily cumbersome to find bugs. We have in the meantime found that it is possible to get a Qt plug-in for Eclipse which would have been of interest to use instead. Rhapsody does also exist as an Eclipse plug-in. To create an Eclipse based development tool chain incorporating Qt and Rhapsody and thereby combining MDD and GUI for development of embedded real-time systems in one tool could be an interesting project.

For the implementation of the architecture itself, we have used Rhapsody. Rhapsody is a very cumbersome program and requires a great insight into how the program works before it really gives a good output for code generation. In this project we have used a large amount of time to get this insight in using Rhapsody. The magnitude of this project could perhaps have been reduced just using Rhapsody for documentation and writing the code manually. The flip-side of the coin is however that when a system is first created in Rhapsody, it is very easy to maintain, add new functionality or change and test the design. Rhapsody produces some code overhead like virtually all code generation systems do.

2.9. Excellence of the project

The perhaps most important part of the design for this project is the layered architecture in which we have divided the design into different packages. Especially the separation of the application into the discrete and continuous packages where the dependency is minimized by use of the façade, observer and command pattern in interchanging information between these two layers.

The application is made very easy to move to another platform by the HW and OS abstraction layers and use of Qt. We have already proved this by the OS abstraction implement for Rhapsody modelling and testing. The design is also very flexible since it would be easy to add a new physiological patient model by use of the strategy in combination with the filter and pipes patterns.

The design is very generic in the way we have designed handling the protocol for the infusion pump. It would be easy to add another external device using the mediator pattern to interact with the patient model. This new device just has to implement the interface specified by the class ProtocolColleague see Figure 20.

The user interface is designed using Qt and the command pattern combined with the state pattern. We have used the architectural pattern Model-View-Controller [6] that is in family with the observer pattern in combination with the way Qt handles the GUI design. It would be possible to add a new view to present a FFT plot of the ECG signal without changing the real-time model. New events can

easy be added for the state machine of the GUI controller and new parameters could be defined setting values in the patient simulator. This is possible due to use of the command pattern.

The tool chain setup we have made enables the developer to continue with an iterative design methodology based on Model Driven Development in Rhapsody. New design patterns can be added to the existing model and verified before generating code to Linux host platform. The GUI can continue to be developed manually in Qt for the discrete package of the application integrating the whole project debugging and testing on the Linux host and cross-compiling to target.

The excel¹⁸ file we have created that contains an RMA scheduling analysis can be used for future extensions to calculate if the scheduling of threads will be possible. This could be the case if the product is extended with new computation of signals then the measurements of WCE must be updated or if configuration parameters for the product are changed. (Sample rate, graph updating rate)

2.10. Suggestion to improvements

Although the project is distinguished in many areas, there are also areas where improvements can be made. One obvious improvement would be to finish the remaining use cases. For now, the system is an application with lot of functionality, but without exploiting it. By implementing the use case 4 and 5 we will achieve a great improvement of the system and expand the opportunities that already functionally exist.

The current system has a problem with a memory leak when killing threads and creating new records. It is obvious that if a "previous" record is not released it will causes memory leaks since the size of the extra memory consumption is the same size as the record file. Therefore, an improvement should be implemented to eliminate these leaks. The use of smart pointers to encapsulate the record pointer would be a solution. By use of smart pointers, the pointer will be encapsulated into a class that will handle deleting the record when it's not used anymore.

Abstraction of the OS could be improved by introducing factory patterns for instantiating all hardware related objects. Currently you must enter and modify all files that use hardware, if you need another abstraction, this could be made easier by centralizing the creation using factory patterns.

It could be interesting to look at some of the many patterns that are not used in the current design. Like use of the adapter pattern to include the legacy "C" code for generation of the EDR signal or see

¹⁸ See CD-rom directory documents and Sapien190_RMA.xls

how the memento pattern could be used to implement an undo function in combination with the command pattern for the discrete part of the system. Other relevant patterns could be: Composite, Template Method, Decorator, Bridge and Flyweight just to mention the ones we think could be relevant.

3. Conclusion

In this project we have achieved to create a patient simulator prototype running on an embedded Linux platform. The development of this prototype is based on the theory from the graduate course "Embedded Real-Time Systems" (TI-IRTS).

The project has incorporated the analysis and design phases of the ROPES spiral micro-cycle model that covers use case requirement specification and domain analysis based on a selection of use cases for an iterative development process. We have combined the ROPES methodology with the basics ideas from Scrum. This process has been an efficient and structured way to focus on "What to do?" and "How to do it?". It has though been hard always to follow this strategy.

In the design phase we have used the 4+1 view model for defining the architectural design of the prototype. We have gained experience in applying many different architectural and design patterns in the final prototype. In this process we have demonstrated how to evaluate and choose design patterns for an embedded real-time system. One of the challenges has been to combine the traditional GoF patterns with patterns targeting the embedded real-time world including concurrency, resource and memory handling. The rate monotonic analysis of the concurrent design we have made can be used in the future to evaluate extension to the product by either adjusting parameters or updating the measurements of WCE that is available in an excel sheet.

The layered architecture and use of patterns between packages has reduced the coupling between layers in the architectural design. This approach makes the product easier to maintain and to continue the development of the individual layers independently. We have achieved this flexibility by use of the design patterns façade, mediator, observer and command. A number of different types of design patterns has been applied and perhaps in some extend "over designed", but it has been fun trying to apply them in a real project.

We have managed to combine Model Driven Development (MDD) by using Rhapsody, GUI development with Qt and cross-compilation to the target in a combined tool chain. This approach makes it possible to maintaining the essential design in the Rhapsody UML model by high level testing and final code-generation to the Linux host integrated with Qt. The GUI and discrete part of the

system can be maintained in Qt integrated with the Rhapsody generated code testing on the host, by abstracting the hardware before the final cross-compilation and deployment to target. This part of the project has perhaps been a little out of scope in relation to the learning objectives for this course, but we have gained valuable knowledge in the field of embedded Linux application development based on Qt and Rhapsody.

The design has plenty of room for extensions to complete the product with all the use cases we have not yet implemented (Like completion of UC#2, UC#4 and UC#5). We have spend most of the time in the inception and elaboration phases of the unified process and therefore work still needs to be done to implement, debug and test a final product.

To conclude, we have created an embedded real-time product based on a number of design patterns and OOA/D methodologies that is very easy to extend and maintain, as extensions will require addition of code and not modification of existing code. We have used an agile development process and complex tools for MDD and GUI development to cut development time.

4. References

1. Gamma, Erich et al. Design Patterns: Elements of Reusable Object-Oriented Software, Addison Wesley (GoF) 1994
2. Douglass, Bruce Powel. Real-Time Design Patterns: Robust Scalable Architecture for Real-Time Systems, Addison Wesley 2002
3. PhysioNet and PhysioBank the research resource for complex physiologic signals.
<http://www.physionet.org/>
4. System/Product architecture document for Sapien 190
<http://code.google.com/p/iirtsf10grp5/downloads/detail?name=Sapien190Spec.doc&can=2&q=>
5. Requirement specification for Sapien 190
<http://code.google.com/p/iirtsf10grp5/downloads/detail?name=Sapien190Spec.doc&can=2&q=>
6. Buschmann, Meunier, Rohnert, Sommerlad, Stal, Pattern-Oriented Software Architectures
7. Project description for Patient Simulator System (PSIMU)
<http://kurser.iha.dk/eit/tiirts/Projekter/PSIMU-project.doc>
8. Larman, Craig. Applying UML and Patterns, An Introduction to Object-Oriented Analysis and Design and Iterative Development. Prentice Hall 2008
9. Shaw, Mary ; Garlan, David. Software Architecture, Perspectives on an Emerging Discipline, Prentice-Hall 1996